

IBM Training

Red Hat OpenShift Application

Development

Luca Floris
IBM Cloud Technical Consultant



Agenda

Day 4

DevOps: continuous delivery

- DevOps and DevSecOps
- WebSphere Liberty on OpenShift
- Transformation Advisor - Migrating WebSphere Applications to OpenShift

Microservices architecture

- Microservices application architecture
- Developing microservices
- Twelve factor applications
- Refactoring monolith applications into microservices

DevOps

What is DevOps?

It is an approach on the journey to lean and agile software delivery that promotes closer collaboration between lines of business, development, and IT operations while removing barriers between your stakeholders, and your customers.

What is DevOps?

Development teams need to design, develop, deliver and run the software as quickly and reliably as possible.

Operations teams need to identify and resolve problems as soon as possible by monitoring, predicting failure, managing the environment and fixing issues.

Benefits of DevOps

Implementing a DevOps practice can add value to your organization through a number of benefits, including the following:

- ✓ Faster code delivery
- ✓ Faster time to market
- ✓ Higher-quality software
- ✓ Improved collaboration between developers and operations
- ✓ Decreased time to resolution for fixing bugs and vulnerabilities
- ✓ A culture that brings business, development, and operations together for improved responsiveness to market demands



DevOps Tools

Define and plan - which focuses on planning DevOps workflows for iterations, release management, and issue tracking.

Code, build, and configure - which focuses on code development and review, source code management, and code merging.

Test - which verifies that the quality of the software release and code are maintained throughout the development process and that the highest quality deploys to production.

Packaging and preproduction - which refers to the activities involved once the release is ready for deployment; it's also called staging or preproduction.

Release, deploy, and orchestration - which is the process of actually releasing software and usually involves change management, release approvals, release automation, schedule orchestration, provisioning, and deploying into production.

Continuous management and configuration - includes continuous configuration automation, configuration management, and infrastructure as code.

Monitoring - reports application performance and helps identify issues impacting the user experience.

DevOps Methodologies

Continuous integration - which is where coding, building, integrating, and testing take place.

Continuous delivery - which includes continuous integration, but mainly focuses on product releases.

Continuous deployment - which focuses on automating releases of projects as soon as possible.

Operate - for conducting the development operations of configuration management and continuous monitoring.

DevOps Principles

Process improvement initiatives to truncate feedback loops to continuously implement needed bug fixes and vulnerability remediation earlier and more cost effectively

Continual experimentation that encourages risk-taking and learning from success and failure, so continuous attempts will lead to future success and mastery

Learners becoming teachers and passing along their acquired knowledge to their colleagues

Using DevOps automation to improve efficiency

Giving continuous feedback to the entire organization

Incentivizing development, test, and deployment teams to collaborate on shared goals

DevOps Strategies

DevOps strategy focuses on the enterprise capability for continuous software delivery that enables customers to seize market opportunities and reduce time to customer feedback

Accelerate the delivery of reliable software

Balance speed, cost, quality, and risk with increased capacity to innovate

Reduce time to customer feedback with improved customer experience

Continuous Integration

Frequently code, build, integrate, and test the work of all developers on a software project at least once a day

DevOps vs Agile

Agile enables developers to deliver their functions every two weeks to respond to changing business needs. DevOps focuses on the operational side of the software development lifecycle, reducing handoffs from developer to operations teams, reducing the time to testing and deploying code, and decreasing errors and downtime of operational systems.

DevSecOps

DevOps benefits security, and with the right automation and operational tools, you can inject security earlier into the development process and increase the security of production code.

DevSecOps

Secure engineering - Products are developed with strong security and comply with appropriate security standards

Secure deployment and operations - Cloud platform and applications are configured and deployed securely, tested for security vulnerabilities, patched, and have their bugs fixed

Separation of duties - Users access what's required for their job duties

Availability and business continuity management - For 99.999 percent availability of infrastructure, runtime components, and management components

Security evaluation and learning - You maintain security functions and properties in code and services as threats evolve and you discover new vulnerabilities

DevSecOps

Benefits

Observability

Traceability

Confidence

Compliance

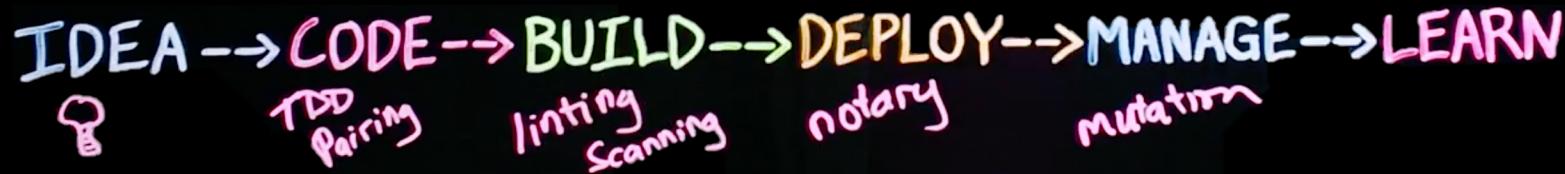
Use Cases

Lack Visibility

Troubled Audits

Unified Governance

Risk



WebSphere Liberty Intro

Why Choose Liberty

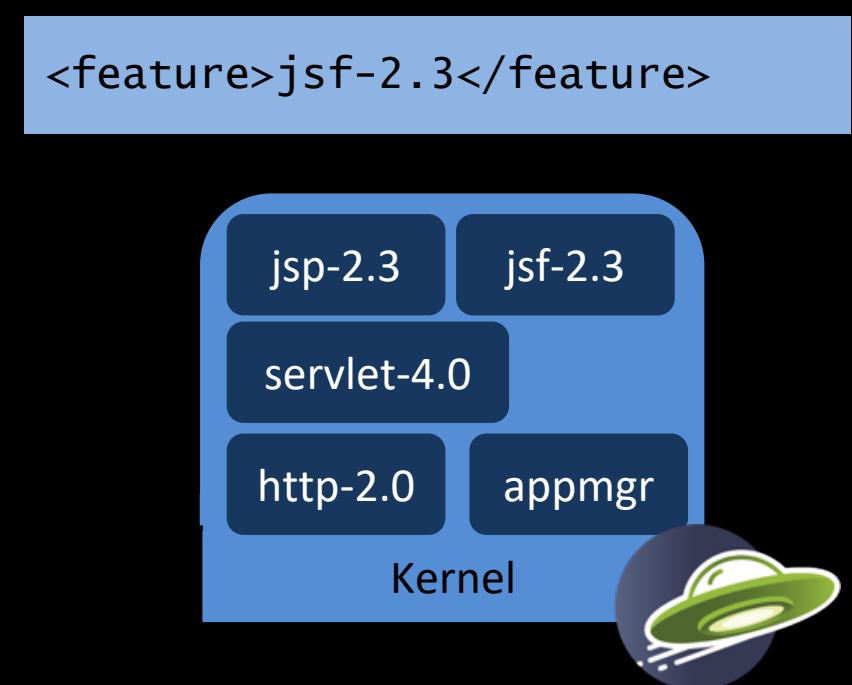
Open Liberty



- Cloud Ready
 - Container optimized
 - Designed for dev/ops
 - Small disk footprint
 - Efficient memory usage
 - Fast startup
 - High throughput
 - Self-Tuned Thread Pool
- Developer friendly
 - Just enough application server
 - Fast inner loop with dev mode
 - Support for industry standard dev tools
 - Jakarta EE, Java EE, MicroProfile APIs
 - Zero Migration
 - Easy install and setup

Just Enough Application Server

You control which features are loaded into each server instance



Periodic Table of Liberty (20.0.0.6)



Periodic Table of Liberty

zOS

ND

Base

Core

Open Liberty

New in 3Q19

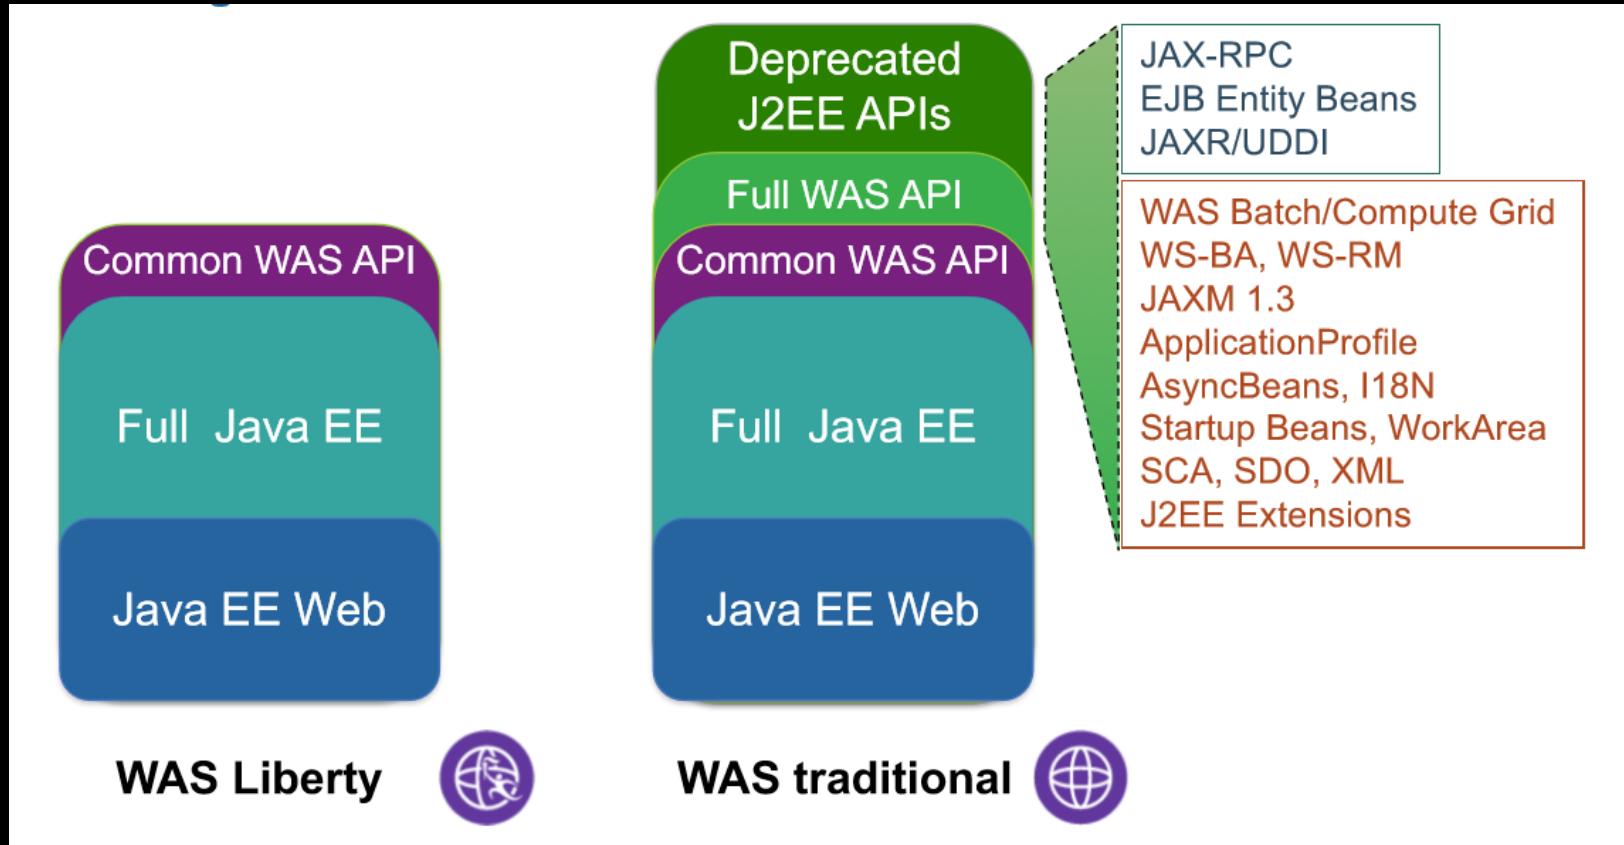
New in 2Q19

New in 1Q19

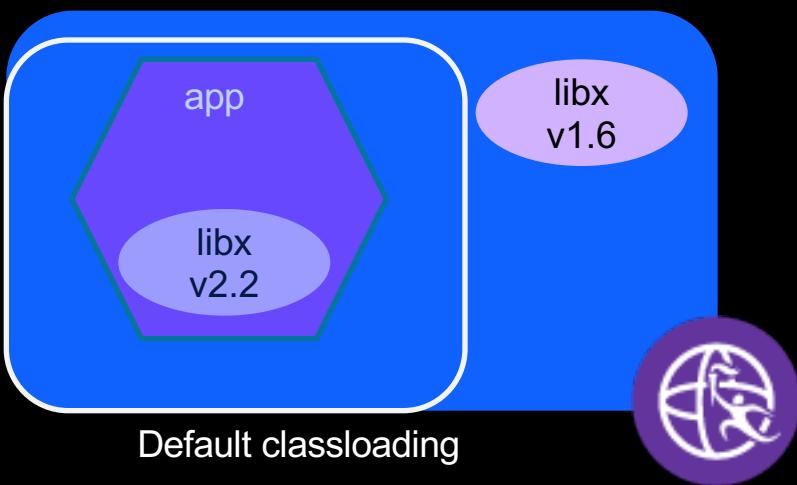
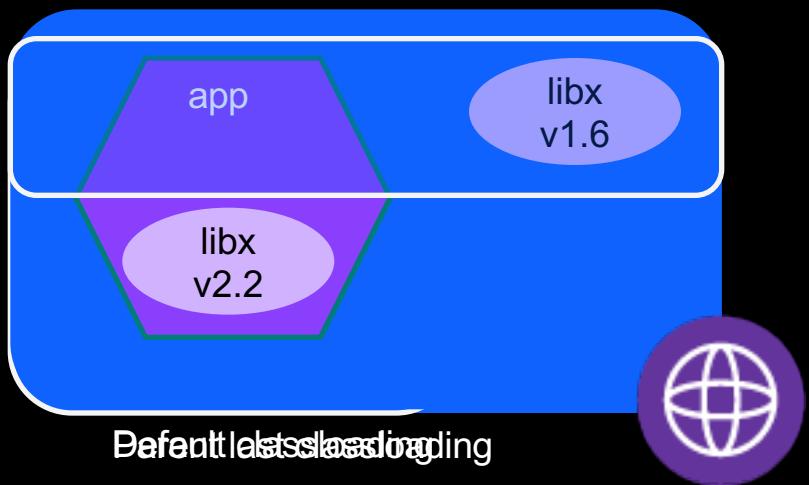
New in 4Q18

		batchSMFLogging-1.0	appClientSupport-1.0	ejbHome-3.2	jacc-1.5	managedBeans-1.0
		collectiveController-1.0	appSecurityClient-1.0	ejbPersistentTimer-3.2	jaxb-2.2	mdb-3.2
clusterMember-1.0		dynamicRouting-1.0	batch-1.0	ejbRemote-3.2	jaxws-2.2	wasJmsClient-2.0
		cloudant-1.0	concurrent-1.0	j2eeManagement-1.1	jca-1.7	webProfile-8.0
		javaee-7.0	ejb-3.2	javaMail-1.6	jms-2.0	wmqJmsClient-2.0
		javaee-8.0				websocket-1.1
		sipServlet-1.0				
		bells-1.0	appSecurity-3.0	jaxrs-2.1	jsonp-1.1	websocket-1.1
		concurrent-1.0	beanValidation-2.0	jaxrsClient-2.1	jsf-2.3	
		javaMail-1.6	cdi-2.0	jdbc-4.2	jsp-2.3	
		jaxb-2.2	ejbLite-3.2	jndi-1.0	managedBeans-1.0	
jdbc-4.3		opentracing-1.3	el-3.0	jpa-2.2	servlet-4.0	
		osgiConsole-1.0	jaspic-1.1	jsonb-1.0	ssl-1.0	
		jpaContainer-2.2				
		springBoot-2.0				
		jsfContainer-2.3				
		webProfile-7.0				
		json-1.0				
		jsonbContainer-1.0				
		jsonpContainer-1.1				
microProfile-3.0			cdi-2.0	jsonb-1.0	mpMetrics-2.0	mpOpenTracing-1.3
			jaxrs-2.1	mpConfig-1.3	mpJwt-1.1	mpRestClient-1.3
			jsonp-1.1	mpFaultTolerance-2.0	mpOpenAPI-1.1	mpHealth-2.0
APIs						

Liberty and WAS API differences



Class Visibility



Liberty Server Configuration

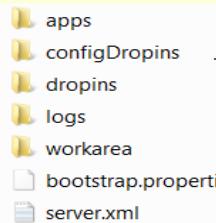
simple but powerful and flexible

```
<server description="new server">  
  
    <!-- Enable features -->  
    <featureManager>  
        <feature>servlet-4.0</feature>  
        <feature>jdbc-4.2</feature>  
    </featureManager>  
    <httpEndpoint id="defaultHttpEndpoint"  
        host="${env.COMPUTERNAME}"  
        httpPort="${httpAdminPort}"  
        httpsPort="${httpsAdminPort}" />  
    <webApplication id="blogapp"  
        location="blogapp.war"/>  
    <include location="${shared.config.dir}/datasource.xml"/>  
</server>
```

You define **which features** you want to use and in **which version**. Features can be **activated on the fly**.

Use **environment variables** or **properties** to **share and re-use** server configurations.

Use **includes** to split the configuration into useful parts for **split of responsibility** and **re-use**.



Use the **configDropins** directory to **override** configurations

Security

- Security by default
 - No remotely accessible ports

```
<httpEndpoint id="defaultHttpEndpoint" host="*"/>
```

- Enable admin, enable security

```
<feature>restConnector-1.0</feature>
<quickStartSecurity userName="admin"
    userPassword="{aes}adSDwijgnb=="/>
```

- Enable ssl using ssl-1.0

```
<feature>ssl-1.0</feature>
<keyStore password="{aes}adSDwijgnb=="/>
```

App Security

- Feature to enable
- Configure security role bindings in server.xml

```
<feature>appSecurity-2.0</feature>
<webApplication location="myweb.war">
    <application-bnd>
        <security-role name="user">
            <group name="myGroup"/>
        </security-role>
    </application-bnd>
</webApplication>
```

- Configure registry

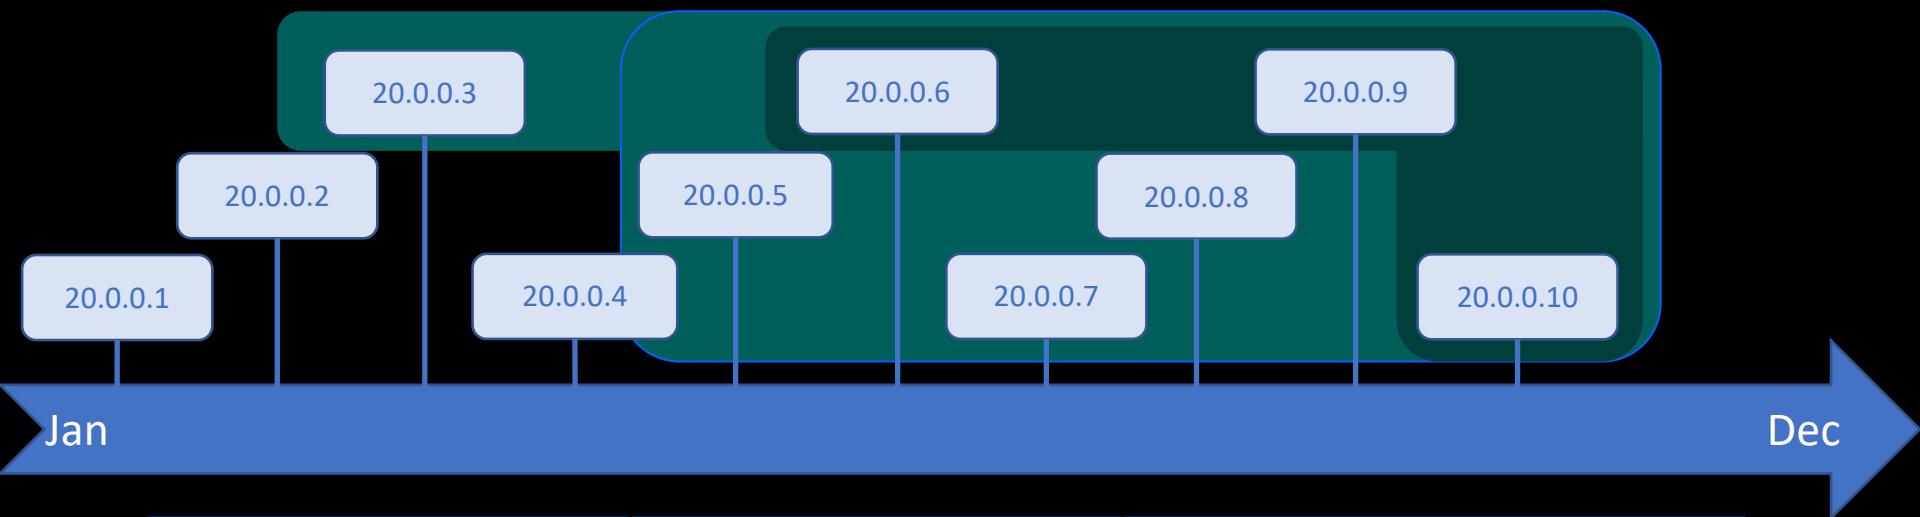
```
<feature>ldapRegistry-3.0</feature>
```

Liberty Zero Migration

- Zero config migration
 - Write once, run forever
- Zero app migration
 - No behavior changes in existing features
 - New behaviors in new features
- Choose your Java
 - Java 14, 11, 8
 - AdoptOpenJDK
 - IBM
 - OpenJDK
 - Oracle



Liberty Release Cadence Example



	All CD releases	CD releases ending .3 .6 .9 .12
Support Provided	5 years	5 years
iFixes	24 weeks	2 years
Proactive Security iFixes	Most recent	Most recent 2

Build

Integrate into build environment

APIs provided in Maven Central

Plugins for Maven and Gradle

Dev mode allows you to develop with any text editor providing hot reload and deployment, on demand testing, and debugger support

```
<plugin>
    <groupId>io.openliberty.tools</groupId>
    <artifactId>liberty-maven-plugin</artifactId>
    <version>3.2</version>
    <configuration>
        <serverName>guideServer</serverName>
    </configuration>
</plugin>
```

The screenshot shows a list of five Open Liberty packages available on Maven Central:

- 5. Open Liberty JavaEE8 Package**
io.openliberty » [openliberty-javae8](#)
Open Liberty JavaEE8 Package
Last Release on Oct 23, 2020
- 6. Open Liberty MicroProfile 3 Package**
io.openliberty » [openliberty-microProfile3](#)
Open Liberty MicroProfile 3 Package
Last Release on Oct 23, 2020
- 7. Open Liberty Kernel Package**
io.openliberty » [openliberty-kernel](#)
Open Liberty Kernel Package
Last Release on Oct 23, 2020
- 8. Open Liberty JavaEE 8 Web Profile Package**
io.openliberty » [openliberty-webProfile8](#)
Open Liberty JavaEE 8 Web Profile Package
Last Release on Oct 23, 2020
- 9. Open Liberty Runtime Package**
io.openliberty » [openliberty-runtime](#)
Open Liberty Runtime Package
Last Release on Oct 23, 2020

Customized Docker containers

Liberty images on Docker Hub

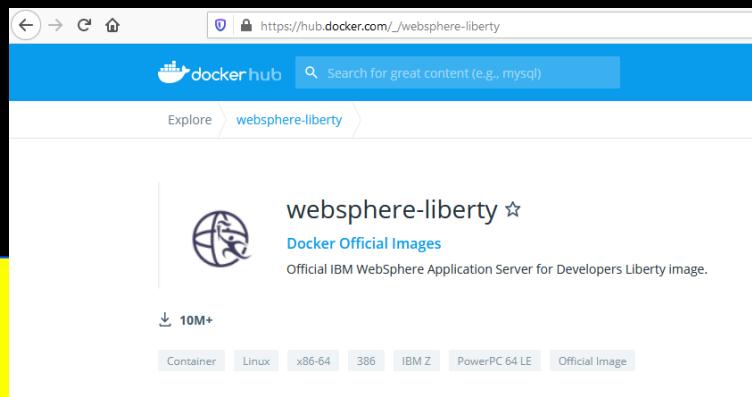
- Liberty containers:

- Kernel, webProfile, full Profile and latest Beta images,
- on Ubuntu or UBI
- with java8 or java11
- Docker files:

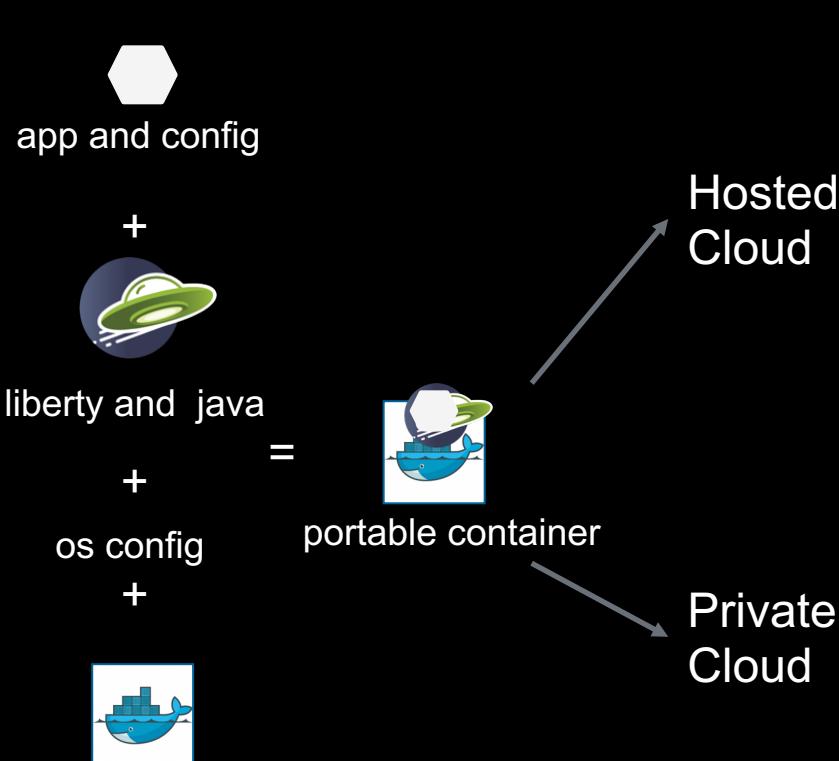
<https://github.com/WASdev/ci.docker>

<https://github.com/OpenLiberty/ci.docker>

```
# FROM openliberty/open-liberty:kernel-java8-openj9-ubi
FROM websphere-liberty:kernel
COPY server.xml /config/
COPY myapp.war /config/apps/
RUN configure.sh
```



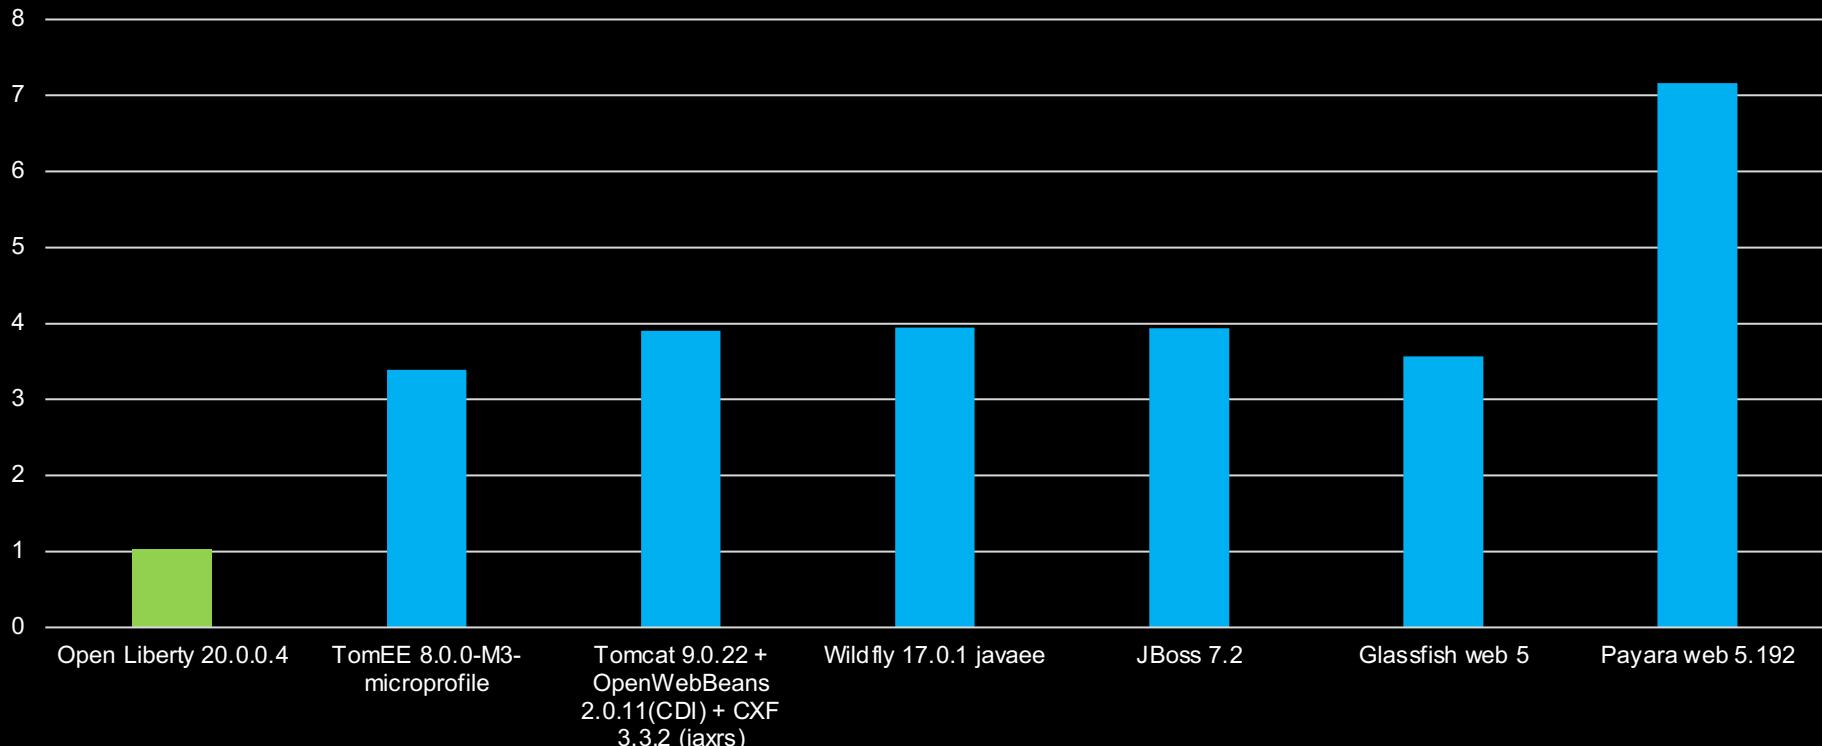
Liberty in Containers



```
FROM open-liberty  
COPY myapp.war /config/dropins/myapp.war
```

Open Liberty startup time comparison (using OpenJ9 JVM)

PingPerf application startup time with OpenJ9 Shared Classes Cache (in seconds)



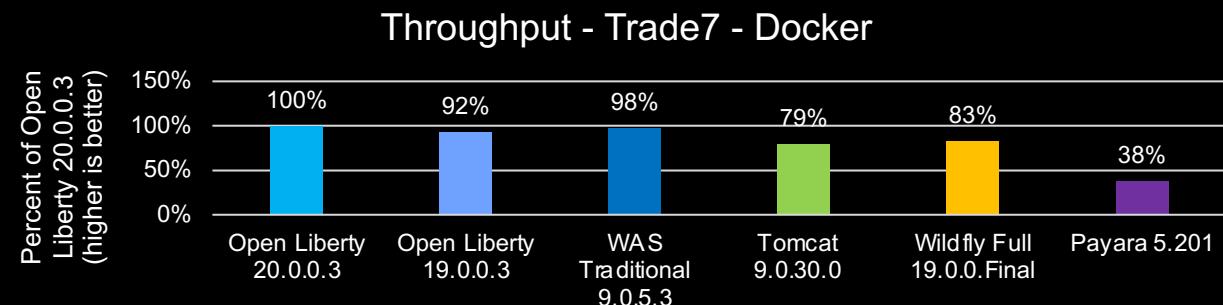
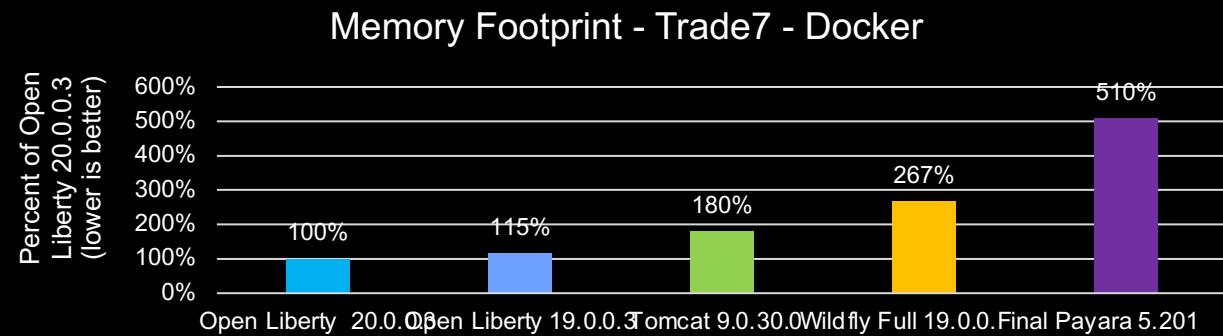
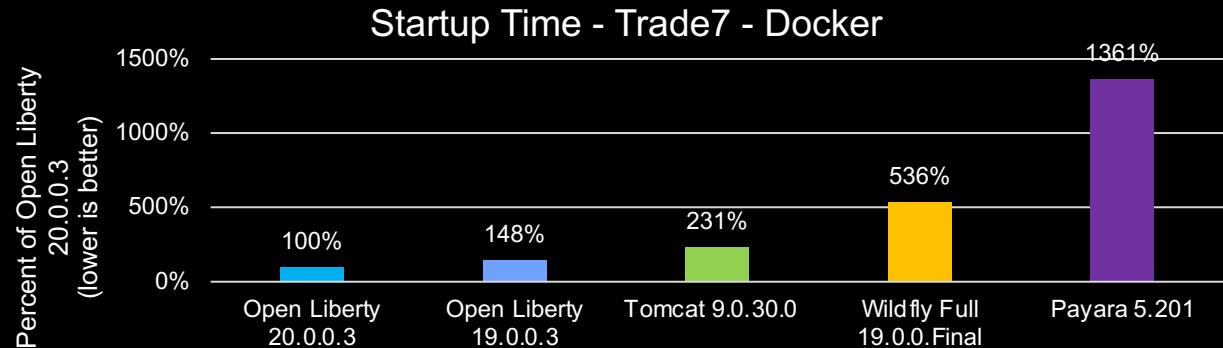
Open Liberty starts faster than other app servers, OpenJ9 starts faster than HotSpot

EE7 Performance (Trade7)

- Comparisons used each application server's Docker image
- Liberty outperforms others on all metrics for EE7 performance (startup time less than half, throughput and memory footprint are much better as well)

System Configuration:

SUT: LinTel – SLES 12.3, Intel(R) Xeon(R) Platinum 8180 CPU @ 2.50GHz, 4 physical cores, 64GB RAM. JDK version distributed with the docker images used for each server instance.



Open Liberty Guides

- Hands-on learning in ~20 minutes
- 44 guides
 - MicroProfile & Jakarta EE
 - Open Shift, Docker, Kubernetes Istio
- Latest Guides
 - *Deploying microservices to an OKD cluster using Minishift*
 - *Deploying microservices to Google Cloud Platform*

The screenshot shows the 'Guides' section of the Open Liberty documentation. At the top, there's a navigation bar with 'Docs', 'Guides', and 'Reference'. Below it, a breadcrumb trail says 'Docs > Guides'. The main title is 'Guides' with a subtitle 'The quickest way to learn all things Open Liberty, and beyond!'. Under this, there's a heading 'Open Liberty Basics - Let's get started' followed by '4 essentials'. Three cards are shown: 'Packaging and deploying applications' (25 minutes), 'Building a web application with Gradle' (15 minutes), and 'Building a web application with Maven' (15 minutes). Below this, there's a heading 'MicroProfile - Developing microservices with ease' followed by '5 essentials'. Three cards are shown: 'Creating a RESTful web service' (30 minutes), 'Injecting dependencies into microservices' (15 minutes), and 'Consuming RESTful services with template interfaces' (20 minutes). At the bottom, there's a card for 'Testing a MicroProfile or Jakarta EE application'.

<https://openliberty.io/guides>

WebSphere on OpenShift

Liberty Operator

Deploy and manage applications running on Open Liberty into [OKD](#) or [OpenShift](#) clusters

Perform Day-2 operations such as gathering traces and dumps

<https://github.com/OpenLiberty/open-liberty-operator/blob/master/doc/user-guide.adoc>



Operator Backed

Open Liberty Application
provided by IBM

Configuration for the deployment
of an Open Liberty application



Operator Backed

Open Liberty Dump
provided by IBM

Day-2 operation for generating
server dumps



Operator Backed

Open Liberty Trace
provided by IBM

Day-2 operation for gathering
server traces

OperatorHub

Discover Operators from the Kubernetes community and Red Hat partners, curated by Operators on your clusters to provide optional add-ons and shared services to your cluster, providing a self-service experience.

All Items

AI/Machine Learning

Application Runtime

Big Data

Cloud Provider

Database

All Items

liberty



Open Liberty Operator
provided by IBM

Deploy and manage applications...

Installed

Operator key features

- Routing - expose your application to external users via a single toggle.
- High Availability - run multiple instances of your application for high availability (static/autoscaling)
- Persistence - enable persistence for your application by specifying storage requirements.
- Serviceability - easily use a single storage for serviceability related operations, such as gathering server traces or dumps.
- Service Binding - easily bind to available services in your cluster.
- Knative - deploy your serverless application on Knative using a single toggle.

Liberty Operator - Application

```
apiVersion: openliberty.io/v1beta1
kind: OpenLibertyApplication
metadata:
  name: my-liberty-app
spec:
  applicationImage: quay.io/my-repo/my-app:1.0
  service:
    type: ClusterIP
    port: 9080
    expose: true
  storage:
    size: 2Gi
    mountPath: "/logs"
```

Liberty Operator - Request server dump

Application needs to have storage for serviceability already configured

CR must be created in the same namespace as the Pod to operate on.

Dump file name will be added to OpenLibertyDump CR status and file will be stored in serviceability folder using format such as
/serviceability/NAMESPACE/POD_NAME/TIMESTAMP.zip

Once the dump has started, the CR can not be re-used to take more dumps. A new CR needs to be created for each server dump.

```
apiVersion: openliberty.io/v1beta1
kind: OpenLibertyDump
metadata:
  name: example-dump
spec:
  podName: Specify_Pod_Name_Here
  include:
    - thread
    - heap
```

To see the status of all dump operations in the current namespace run:

```
oc get oldump -o wide
```

Liberty Operator - Request server traces

Application needs to have storage for serviceability already configured

CR must be created in the same namespace as the Pod to operate on.

Generated trace files, along with *messages.log* files, will be in the folder using format
/serviceability/NAMESPACE/POD_NAME/

To stop the trace:

- Set the disable parameter to true
- Or delete the CR

```
apiVersion: openliberty.io/v1beta1
kind: OpenLibertyTrace
metadata:
  name: example-trace
spec:
  podName: Specify_Pod_Name_Here
  traceSpecification: "*=info:com.ibm.ws.webcontainer*=all"
  maxFileSize: 20
  maxFiles: 5
```

To see the status of all trace operations in the current namespace run:

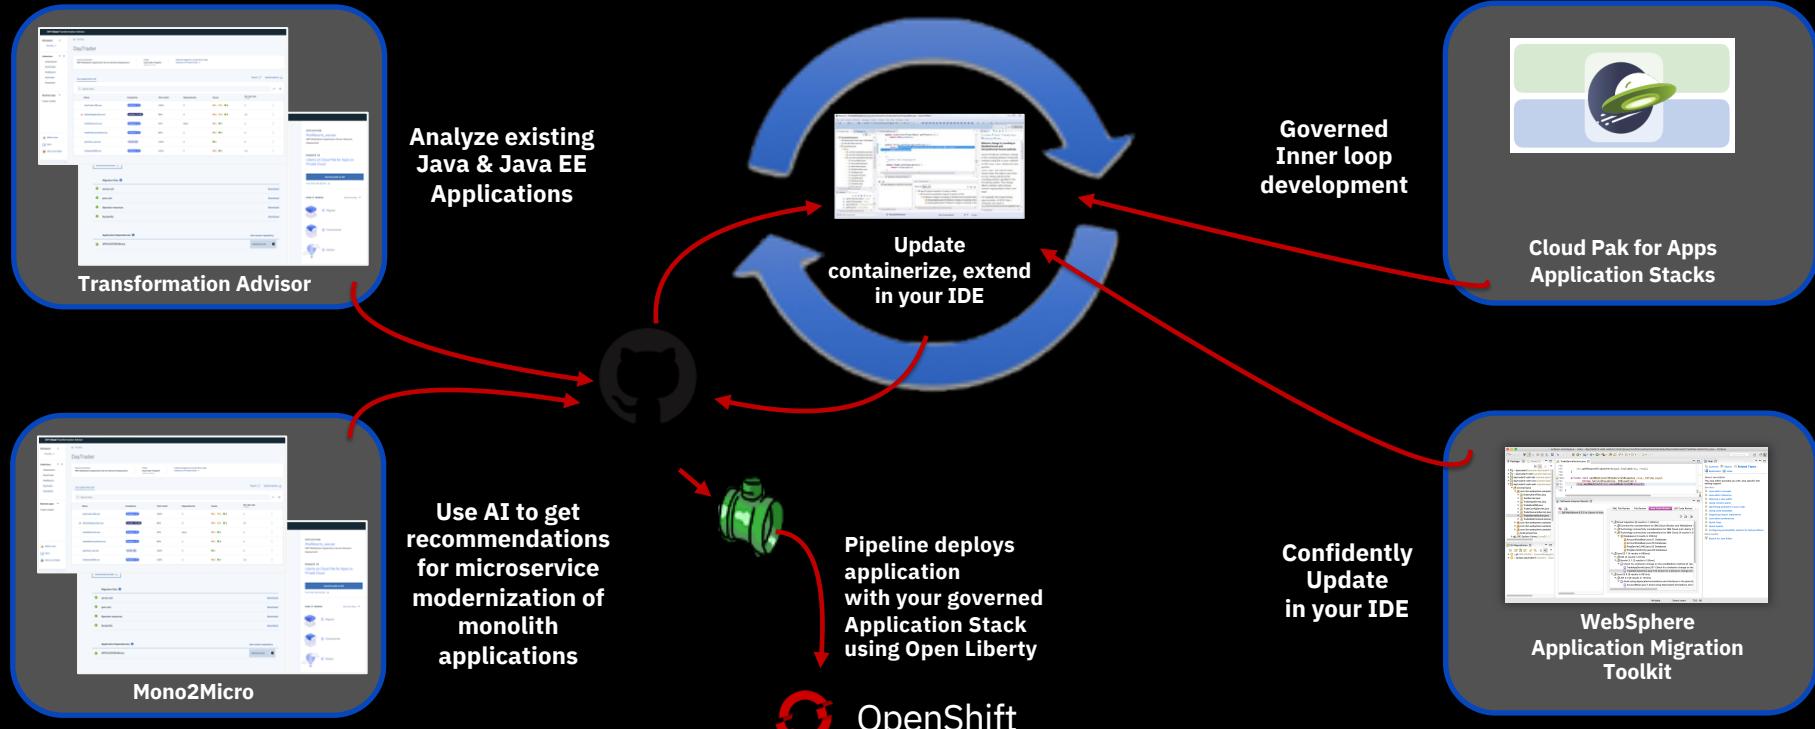
```
oc get oltrace -o wide
```

Resources

- Open Liberty: <https://www.openliberty.io/>
- Eclipse MicroProfile: <https://microprofile.io/>
- Jakarta EE: <https://jakarta.ee/>
- Liberty advantage: <https://www.ibm.com/downloads/cas/NVY3KY4E>
- Open Liberty Guides: <https://openliberty.io/guides>
- Why Liberty is the best Java runtime for the Cloud <https://developer.ibm.com/wasdev/docs/liberty-profile-best-java-runtime-cloud/>
- WebSphere Application Server V8.5 Administration and Configuration Guide for Liberty Profile (Redbook) <http://www.redbooks.ibm.com/abstracts/sg248170.html?Open>
- Liberty videos: https://www.ibm.com/support/knowledgecenter/SSAW57_liberty/com.ibm.websphere.wlp.nd.multiplatform.doc/ae/covr_media.html
- Java support dates <http://www.ibm.com/developerworks/java/jdk/lifecycle>
- Single Stream Continuous Delivery <https://www-01.ibm.com/support/docview.wss?uid=ibm10869798>
- WebSphere Migration Knowledge Collection: Planning and Resources <https://www-01.ibm.com/support/docview.wss?uid=swg27008724>
- IBM Transformation Advisor <http://ibm.biz/cloudta>
- WebSphere Binary Migration Toolkit: <http://ibm.biz/WAMT4AppBinaries>

Modernization tools to accelerate the Java process end-to-end

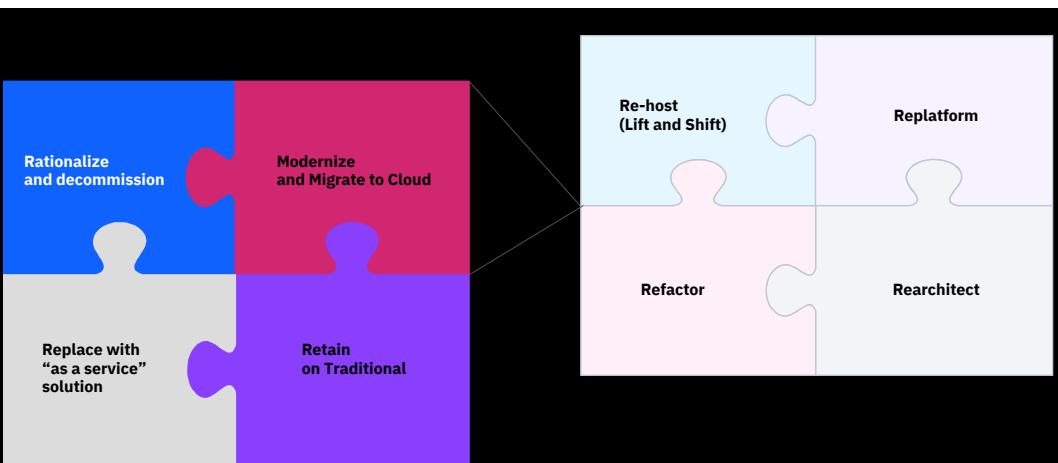
*One development experience for both cloud-native and application modernization
lowering the on-ramp to cloud development, test, deployment for legacy app developers*



The optimum modernization depends on the workload needs

IBM Cloud Transformation Advisor

accelerates that journey by quickly discovering and analyzing on-premise Java EE and/or messaging workloads in the enterprise to help in determining and executing the optimum modernization steps for each.



Runtimes currently analyzed:

Java EE



WebSphere Application Server



Oracle WebLogic



Red Hat JBoss

Apache Tomcat



Integration



IBM MQ

App Connect Enterprise (IIB)

Open Source SDK for other extenders:

<https://github.com/IBM/transformation-advisor-sdk>

Transformation Advisor (TA) includes

Containerized runtime target analyses
(Liberty, WebSphere Base, ACE, MQ)

- Modernization complexity assessment
- Issue identification and severities
- Resolution cost assessments

Holistic views of related components

- Dependency mappings
(shared libs, middleware relationships)
- Business Applications with cumulative issue summary and resolution costs

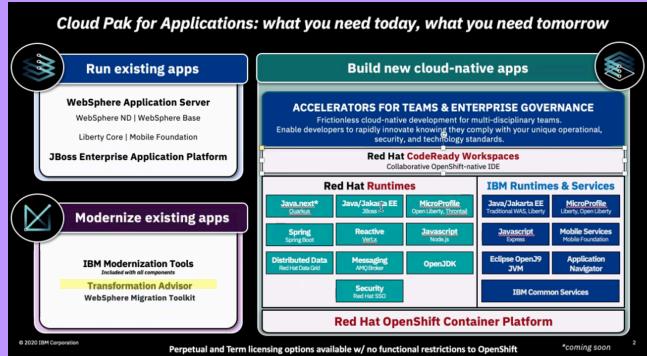
Detailed Reports

- Issue identification and remediation
- Workload inventories
- Dependency identification

Application-specific customized migration plans and deployment artifacts

- Configuration for containerized runtime
- Build and deployment files
- Integrated with Accelerator for Teams

Get Transformation Advisor

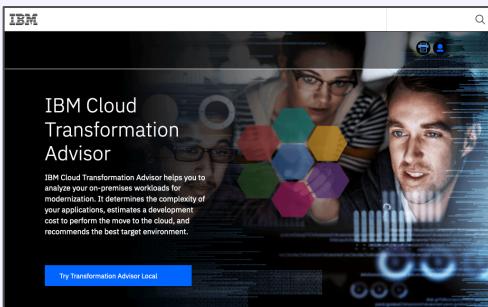


TA is available in
IBM Cloud Pak for Apps
and
IBM Cloud Pak for Integration
(v2020.2.1 – 2Q2020)

The screenshot shows the IBM Cloud Pak for Integration interface. It displays a dashboard with various service icons such as API Management, Data Integration, and Application Integration. Below the dashboard, there are sections for 'IBM Runtimes & Services' (Java/Jakarta EE, MicroProfile, JavaScript, Mobile Services, Eclipse OpenJ9 VM, Application Navigator) and 'IBM Common Services'. A note at the bottom states 'IBM Cloud Pak for Integration Multi-Cloud, Secure, Enterprise-proven Platform'.

- Create all types of Integrations quickly using **simple** tools following modern development practice
- Manage secure integration deployments across any cloud leveraging cloud-native practices & architecture, consistent operational services.
- Re-use integration assets, skills to speed up development and time-to-market maintaining the right governance

[TA 2.1.1 Operator Tile – available on RedHat Operator Hub](#)



<http://ibm.biz/cloudta>

A Trial version of TA is available to run locally in Docker containers:

- MacOS
- Linux
- Windows

Run Transformation Advisor

Open TA

Download the appropriate Data Collector

Run Data Collection and
Upload Resulting Zipfile (if necessary)

View Remediation Assessment to
Target Containerized Runtime(s)

Welcome to Transformation Advisor

We know that modernizing your existing middleware deployments can take you into unfamiliar territory. IBM Cloud Transformation Advisor will help you take your first steps toward getting you onto Cloud Pak for Apps.

Tired of this top section? Hide it

Let's get started.
What's a workspace?

Add a new workspace

SampleData a minute ago

Most recently accessed (2 total)

Join our community to ask questions, give feedback, or just say hello.

Join the Slack community Transformation Advisor docs What's new

IBM Cloud Transformation Advisor

Home / SampleData / NewDataCollector

Profile: WL_OI_latest_DC Version: 9.0.0.0

Migration target: Cloud Pak for Apps Compatible Liberty runtimes

Upload options

Java applications (13)

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
applicationPetstore_war_ear	Open Liberty	Phenomenal	1	▲ 2 ● 2	2.5
applicationPetstore_war_ear	WebSphere Liberty	Simple	1	● 2	0
DefaultApplication.ear	Open Liberty	Complex	3	▲ 4 □ 1 ● 4	14
ibtApp.ear	Open Liberty	Phenomenal	2	▲ 2 □ 1 ● 2	3
modresources-1_0_war.ear	Open Liberty	Simple	None	● 1	0
PlentyByWebSphereEJB.ear	Open Liberty	Modest	3	▲ 1 ● 5	0.5
query.ear	Open Liberty	Phenomenal	2	▲ 2 □ 1 ● 3	3
ReportViewerDeployment.ear	Open Liberty	Phenomenal	3	▲ 2 □ 1 ● 4	12.5
ReportViewerDeployment.ear	WebSphere Liberty	Modest	3	▲ 5 □ 1 ● 8	8
Tax.ear	Open Liberty	Simple	None	● 0	0

Items per page: 10 1-10 of 13 items

IB Cloud Transformation Advisor

Home / SampleData / NewDataCollector

Sessions environment: IBM WebSphere Application Server

Profile: WL_OI_latest_DC Version: 9.0.0.0

Migration target: Cloud Pak for Apps Compatible Liberty runtimes

Java applications (13)

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost
applicationPetstore_war_ear	Open Liberty	Phenomenal	1	▲ 2 ● 2	2.5
applicationPetstore_war_ear	WebSphere Liberty	Simple	1	● 2	0
DefaultApplication.ear	Open Liberty	Complex	3	▲ 4 □ 1 ● 4	14
ibtApp.ear	Open Liberty	Phenomenal	2	▲ 2 □ 1 ● 2	3
modresources-1_0_war.ear	Open Liberty	Simple	None	● 1	0
PlentyByWebSphereEJB.ear	Open Liberty	Modest	3	▲ 1 ● 5	0.5
query.ear	Open Liberty	Phenomenal	2	▲ 2 □ 1 ● 3	3
ReportViewerDeployment.ear	Open Liberty	Phenomenal	3	▲ 2 □ 1 ● 4	12.5
ReportViewerDeployment.ear	WebSphere Liberty	Modest	3	▲ 5 □ 1 ● 8	8
Tax.ear	Open Liberty	Simple	None	● 0	0

Upload options

Analyze On-Premise Workloads

Java EE

WebSphere
(including shared libraries)

Tomcat

WebLogic

JBoss

MQ

IIB/ACE

The screenshot displays five separate workspace panels from the IBM Cloud Transformation Advisor, each showing a list of Java applications and their migration details:

- FinancialProfile:** Shows 10 Java applications.
- Tomcat:** Shows 20 Java applications, including Apache Tomcat, JBoss, and FinancialProfile.
- WebLogic:** Shows 2 Java applications, including Oracle WebLogic.
- ACEv11:** Shows 1 Java application, ModResorts-JBoss.
- JBoss:** Shows 1 Java application, ModResorts-JBoss.

Each panel includes navigation links like "Home / SampleData / [Profile]", "Upload options", and "Migration target on Cloud Pak for Apps". The ACEv11 panel also shows integration server details and a table of issues and total effort.

Complexity Assessment

Classification regarding the difficult of the types of remediation required to move the workload to the desired target runtime environment

Complexity	Java EE	IBM MQ	IIB/ACE
Simple	No code changes needed	DNS reconfiguration required	Admin change required
Moderate	Code changes required	Cluster reconfiguration, changing custom logic (e.g. Exits)	Development change required
Complex	Incompatible technologies; external dependencies	Client Authentication reconfiguration	Difficult develop task or alternate technology required

The image displays three separate screenshots of the IBM Cloud Transformation Advisor interface, each showing a different migration profile:

- FinancialProfile:** Shows a Java application named "HomeLoan_Mortgage_war.war" with a complexity level of "Simple".
- Messaging:** Shows a queue manager with a complexity level of "Simple".
- ACEv11:** Shows an integration server named "server1" with a complexity level of "Complex".

Each screenshot includes navigation elements like "Upload options", search bars, and tabs for "Java applications", "Queue manager", and "Integration server". The bottom of the interface shows pagination controls.

Issues and Severity

Indication of the immediacy and likelihood of the identified issue needing to be remediated before modernizing to the target runtime

Severity	All Assessments
● Suggested	No immediate action required
■ Potential	Investigation required to determine if action required
▲ Critical	Issue must be addressed

The screenshot shows three separate assessments in the IBM Cloud Transformation Advisor:

- FinancialProfile:** Shows Java applications (10) and Shared libraries (1). A modal is open over the main view, highlighting the "Issues" column in the Java applications table. The table includes columns: Name, Migration Target, Complexity, Dependencies, Issues, and Estimated dev cost. Issues are color-coded: green (2), yellow (1), and red (1).
- Messaging:** Shows Source environment (IBM MQ), Installation (MQInstallation01), and Preferred migration (MQ on OpenShift). It lists messaging components like Mortgage_Calculator, Mortgage_InterestRates, and Mortgage_Processing.
- ACEv11:** Shows Source environment (App Connect Enterprise), Node (btmnode), and Preferred migration (ACE on OpenShift). It lists integration servers (server1, server2, server3) and their associated issues. The table includes columns: Integration server, Issues, and Total effort (in days). Issues are color-coded: red (2), yellow (2), and green (2).

Resolution Cost Estimates

Estimate of effort required for remediation of issues.

Java estimates are for development efforts, based on actual services engagement experiences resolving those issues.

IBM Cloud Transformation Advisor

Workspace SampleData

Collections Tomcat JBoss FinancialProfile ACEv11 1 more

Business apps 0 apps created

What's new Docs Join us on Slack

Items per page: 10 1-3 of 3 items

ACEv11

Source environment App Connect Enterprise Node bthnode Preferred migration ACE on OpenShift middlewareVersion:v10.0

Upload options

Integration server

Server	Complexity	Issues	Total effort in days
server1	Complex	▲ 2 □ 1 ○ 2	21
server2	Simple	○ 2 □ 2 ○ 2	8
server3	Simple	○ 3 □ 2 ○ 2	12

Search integration server

Upload options

IBM Cloud Transformation Advisor

Workspace SampleData

Collections Tomcat JBoss FinancialProfile WebLogic ACEv11 1 more

Business apps 0 apps created

What's new Docs Join us on Slack

Items per page: 10 1-2 of 2 items

Messaging

Source environment IBM MQ Installation MQInstallation01 Version: 9.0.2.0 Preferred migration MQ on OpenShift

Queue manager

Queue Manager	Complexity	Issues	Total effort in days
ECB.RATES	Complex	○ 2 □ 2 ○ 2	5.5
MORTGAGES	Simple	○ 2 ○ 1	0

Breakdown Complexity of applications that use this queue manager

Complexity	Count
Simple	1
Moderate	1
Complex	1

Used by Applications that use this queue manager

- Mortgage_Calculator01.ear
- Mortgage_InterestRates_war.war
- Mortgage_Processing_war.war

Upload options

IBM Cloud Transformation Advisor

Home / SampleData / FinancialProfile

Profile FinancialProfile01 Version: 8.5.1.4

Migration target on Cloud Pak for Apps Compatible Liberty runtimes

FinancialProfile

Source environment IBM WebSphere Application Server Network Deployment

Java applications (10) Shared libraries (1)

Java applications

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
HomeLoan_Mortgage_war.war	Open Liberty	Simple	1	○ 2	0
loyalty-level-1_0_war.ear	Open Liberty	Moderate	2	○ 1 □ 4	1
▲ Mortgage_Calculator01.ear	Open Liberty	Complex	4	▲ 3 □ 1 ○ 5	13
Mortgage_Calculator02_war.war	Open Liberty	Simple	1	○ 2	0
				○ 6	0
				○ 1 □ 1	2

Search java applications

Upload options

Dependencies and Relationships

Applications using Shared Libraries

Applications and Queue Managers

IBM Cloud Transformation Advisor

Workspace SampleData + Home / SampleData / FinProfile FinProfile

Java applications (10) Shared libraries (5)

Shared libraries

Java applications (10) Shared libraries (5)

Java applications

Search java applications

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
HomeLoan_Mortgage_war.ear	Open Liberty	Simple	1	0 0 2 0	0
loyalty-level-1_0_war.ear	Open Liberty	Moderate	2	0 1 4 1	0
Mortgage_Calculator01.ear	Open Liberty	Complex	4	3 1 5 13	0
Mortgage_Calculator02_war.ear	Open Liberty	Simple	1	0 0 2 0	0

Breakdown

Complexity of queue managers used by this application

1 Queue mgr	0 Simple	0 Moderate	1 Complex
-------------	----------	------------	-----------

IBM MQ queue managers list

This application has the following IBM MQ queue manager

ECB.RATES

Complexity Dependencies Issues Estimated dev cost in days

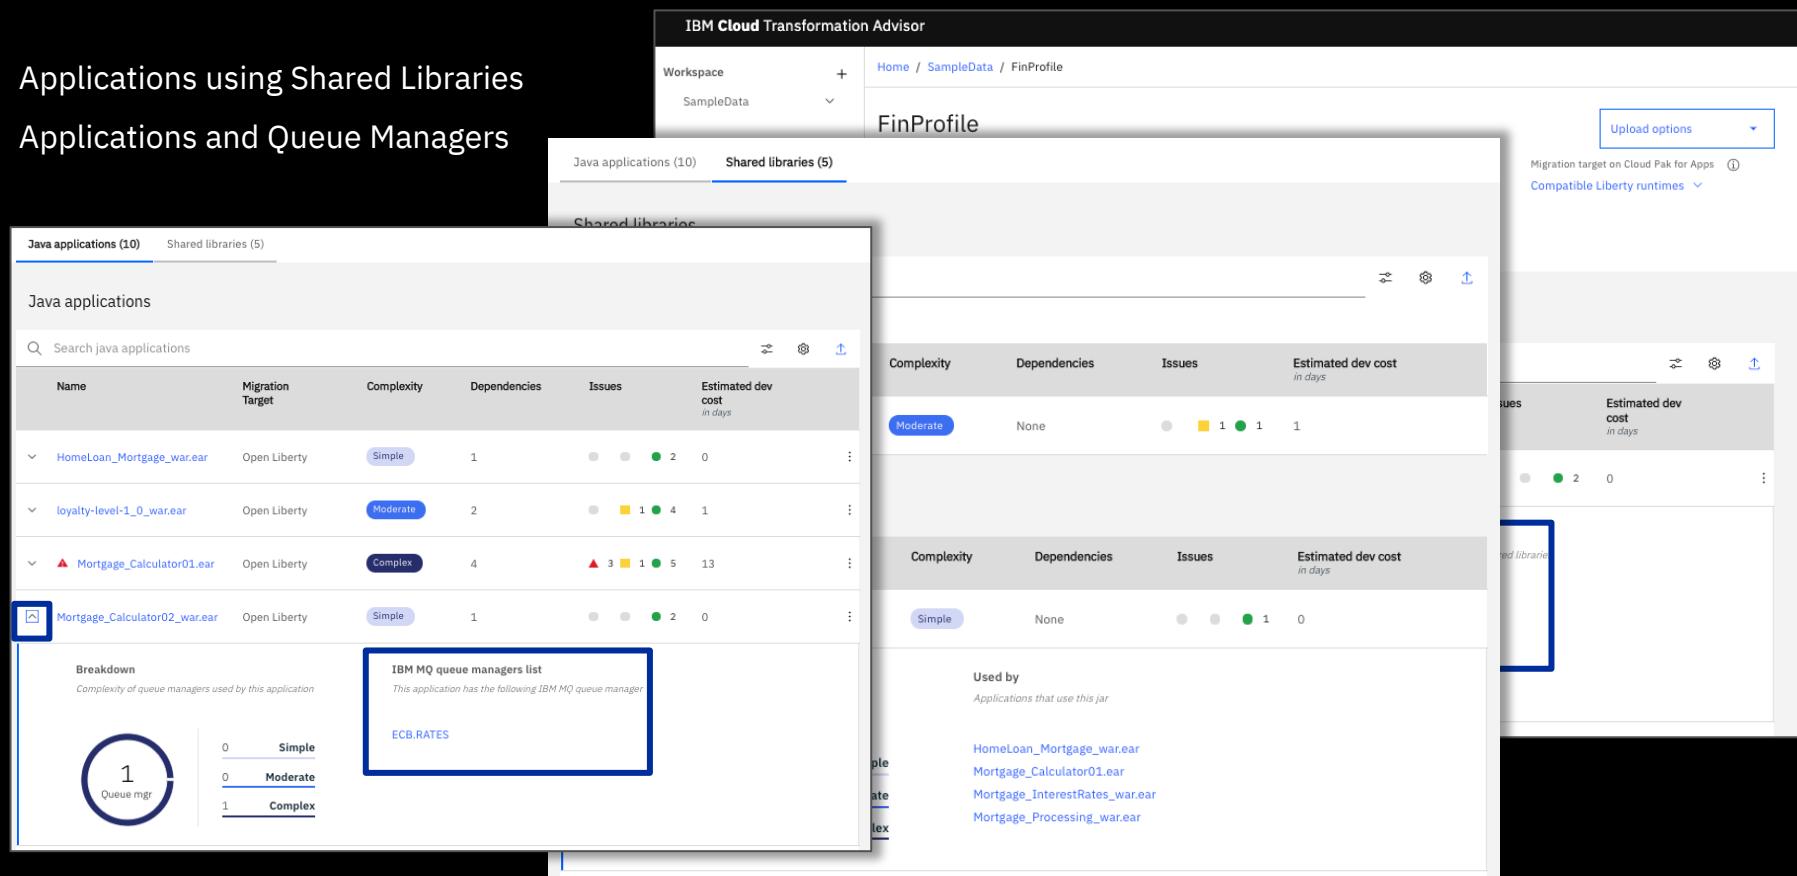
Moderate None 0 1 1 1

Simple None 0 0 1 0

Used by

Applications that use this jar

HomeLoan_Mortgage_war.ear
Mortgage_Calculator01.ear
Mortgage_InterestRates_war.ear
Mortgage_Processing_war.ear



Business Applications

Create Business App from applications in one or more collections

View cumulative issues, development costs and complexity assessment

Includes JARs used by applications from configured shared libraries without double-counting issues/dev cost

The screenshot shows the IBM Cloud Transformation Advisor interface. On the left, there's a sidebar with sections for Workspace (SampleData), Collections (NewDataCollector, SVTApps, SharedLibs, MQ, WAS), and Business apps (MortgageApp selected). The main area is titled "Create business application" and shows a list of components: NewDataCollector, WL_OL_latest_DC, WAS, and FinancialProfile01. It includes filters for complexity (Simple, Moderate) and migration targets (Cloud Pak for Apps, Compatible Liberty runtimes). Below this is a search bar for "Add java applications". To the right, there's a panel for "Business application name" set to "MortgageApp" with options for Java applications (4) and Shared libraries (5). A summary box shows Complexity (Moderate), Issues (1 red, 1 yellow, 18 green), and Estimated dev. cost (3 days). At the bottom, a table lists four Java applications with columns for Name, Migration Target, Complexity, Dependencies, Issues, and Estimated dev cost. The "Mortgage_Processing_war.ear" application has a complexity of Moderate, while others are Simple. The "Mortgage_Processing_war.ear" application also has a higher number of issues (1 red, 1 yellow, 1 green) compared to the others (mostly 2 green). The interface has a "Create" button at the bottom right.

Migration Plans Available for Java Applications

IBM Cloud Transformation Advisor

Home / SampleData / FinProfile

FinProfile

Source environment: IBM WebSphere Application Server
Profile: FinancialProfile01
Version: 8.5.5.14
Migration target on Cloud Pak for Apps
Compatible Liberty runtimes

Upload options

Java applications (10) Shared libraries (5)

Java applications

Search java applications

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
HomeLoan_Mortgage_war.ear	Open Liberty	Simple	1	1 1 2 0	0
loyalty-level-1_0_war.ear	Open Liberty	Moderate	2	1 4 1 1	1
Mortgage_Calculator01.ear	Open Liberty	Complex	4	3 1 5 2	2
Mortgage_Calculator02_war.ear	Open Liberty	Simple	1	1 1 2 0	0
Mortgage_InterestRates_war.ear	Open Liberty	Simple	3	1 6 1 0	0
Mortgage_Processing_war.ear	Open Liberty	Moderate	None	1 1 1 2	2
notification-ear.ear	Open Liberty	Simple	2	1 4 0 0	0
portfolio-1_0_war.ear	Open Liberty	Moderate	3	1 7 1 0.5	0.5

Add to existing business app
Create as a business app
View migration plan

What's new
Docs
Join us on Slack

Open Liberty or WebSphere Liberty option

NEW!

New TA 2.1.0 Data Collector option for applications using Java EE6 versions of:

JPA

JAX-RS

or CDI

The screenshot shows the IBM Cloud Transformation Advisor interface. On the left, there's a sidebar with sections for JPA, JAX-RS, and CDI. The main workspace is named "SampleData". A modal window titled "LibertyOptions" is open, showing a list of Java applications (11) under "Java applications". One application, "ReportViewerDeployment.ear", is highlighted with a red border. The "Migration Target" for this application is set to "WebSphere Liberty". A tooltip on the "Migration target" dropdown says: "Target on Cloud Pak for Apps" and "Compatible Liberty runtimes". Another modal window titled "PreviousDataCollector" is also visible, with a message about a new version of the data collector available.

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
WSBSample.ear	Open Liberty	Moderate	1	▲ 1 ● 3 ○ 2	2
UploadServlet30_war.ear	Open Liberty	Simple	None	● ● ● 0	0
transformationAdv_ear.ear	Open Liberty	Simple	None	● ● ○ 1 0	0
Ta.ear	Open Liberty	Simple	None	● ● ● 0	0
ReportViewerDeployment.ear	WebSphere Liberty	Moderate	3	▲ 5 ○ 1 ● 8 ○ 8	8
ReportViewerDeployment.ear	Open Liberty	Moderate	3	▲ 7 ○ 1 ● 9 ○ 12.5	12.5

Java Modernization Target Deployments

Target Deployments:

Open Liberty

WebSphere Liberty

WebSphere Base

IBM Cloud Transformation Advisor

Workspace + Home / SampleData / NewDataCollector

IBM Cloud Transformation Advisor

SampleData + Home / SampleData / NewDataCollector

NewDataCollector

Upload options

Upload options

Upload options

Upload options

IBM Cloud Transformation Advisor

Workspaces + Home / SampleData / NewDataCollector

NewDataCollector

Source environment IBM WebSphere Application Server ... Profile WL_OI_collection_old_data Version: 8.5.5.11 Migration target on Cloud Pak for Apps WebSphere traditional Estimated dev cost

Java applications (6)

Java applications

Search java applications

Name	Migration Target	Complexity	Dependencies	Issues	Estimated dev cost in days
activitysession.ear	WebSphere traditional	Moderate	4	2 ● 5 1	...
DayTrader EE6.ear	WebSphere traditional	Simple	3	10 0	...
DefaultApplication.ear	WebSphere traditional	Moderate	3	2 ● 4 1	...
ivtApp.ear	WebSphere traditional	Moderate	2	1 ● 3 1	...
PlantsByWebSphereEE6.ear	WebSphere traditional	Simple	3	7 0	...
query.ear	WebSphere traditional	Moderate	2	1 ● 3 1	...

What's new Docs Join us on Slack

What's new Docs Join us on Slack

The screenshot shows the IBM Cloud Transformation Advisor interface. It displays a hierarchy of workspaces and collections, with the 'NewDataCollector' collection selected. The main view shows Java applications (6) listed in a table. Each application row includes columns for Name, Migration Target (e.g., WebSphere traditional), Complexity (e.g., Moderate, Simple), Dependencies, Issues (represented by a bar chart), and Estimated dev cost in days. The 'PlantsByWebSphereEE6.ear' application is highlighted with a blue border. The interface has a dark theme with light-colored cards for each workspace and collection.

Liberty Migration Plans

Source or binary projects

Custom configuration and deployment artifacts created

Full integration with IBM Cloud Pak for Apps Accelerator for Teams

Migration bundle pushed to git

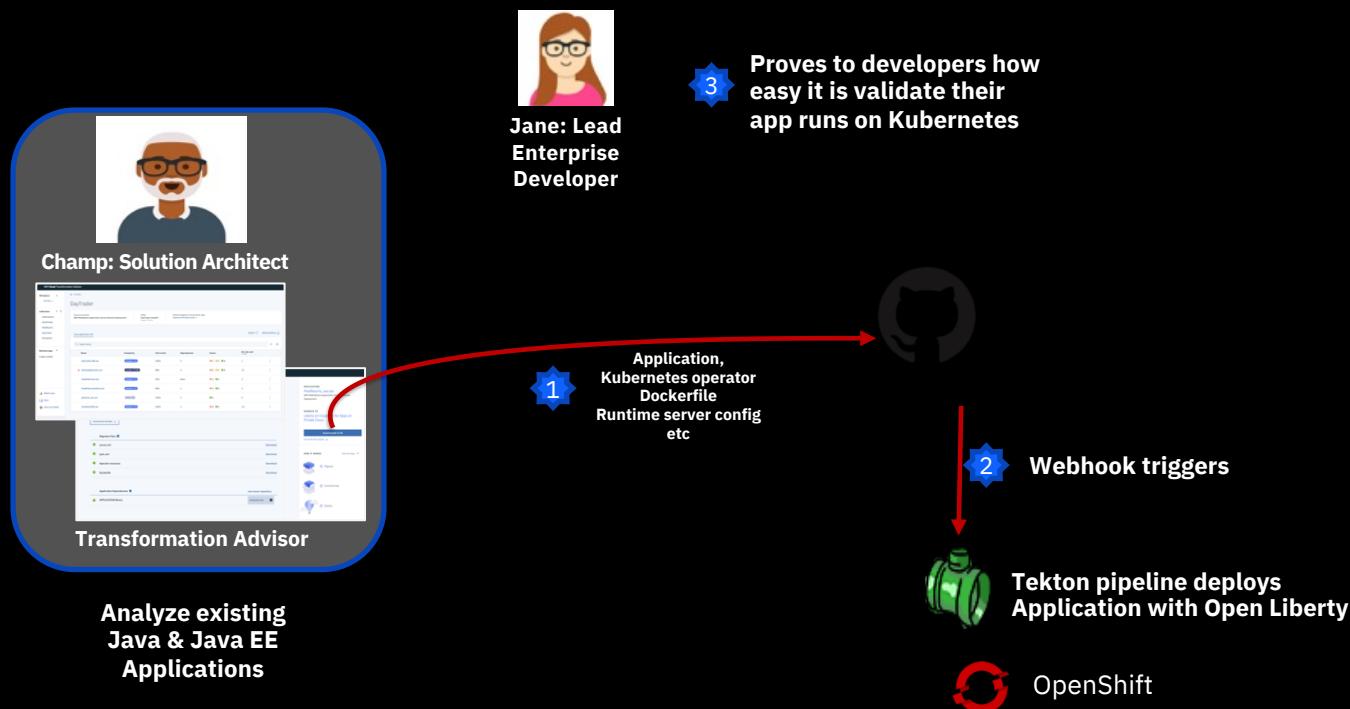
The screenshot shows the IBM Cloud Transformation Advisor interface. At the top, it displays the workspace "SampleData" and the application "NewDataCollector". On the right, there are "Upload options" and a "Send to Git" button highlighted with a blue border.

The main area shows the "Migration bundle" page for the application "ReportViewerDeployment.ear". It includes fields for "Application name" (ReportViewerDeployment.ear), "Source environment" (IBM WebSphere Application Server), and "Migration target" (IBM Cloud Pak for Applications: Open Liberty). A sidebar on the right shows a tree structure with nodes like "Business app plan" and "Business app plan" highlighted with a blue border.

A large callout box highlights the "Use Accelerator for Teams" section, which contains a toggle switch labeled "Use" and a note about using Accelerator for Teams or manually setting up a Git repository. Below this, a "Migration Files" section lists generated files: ".appody-config.yaml", "app-deploy.yaml", "pom.xml", and "server.xml".

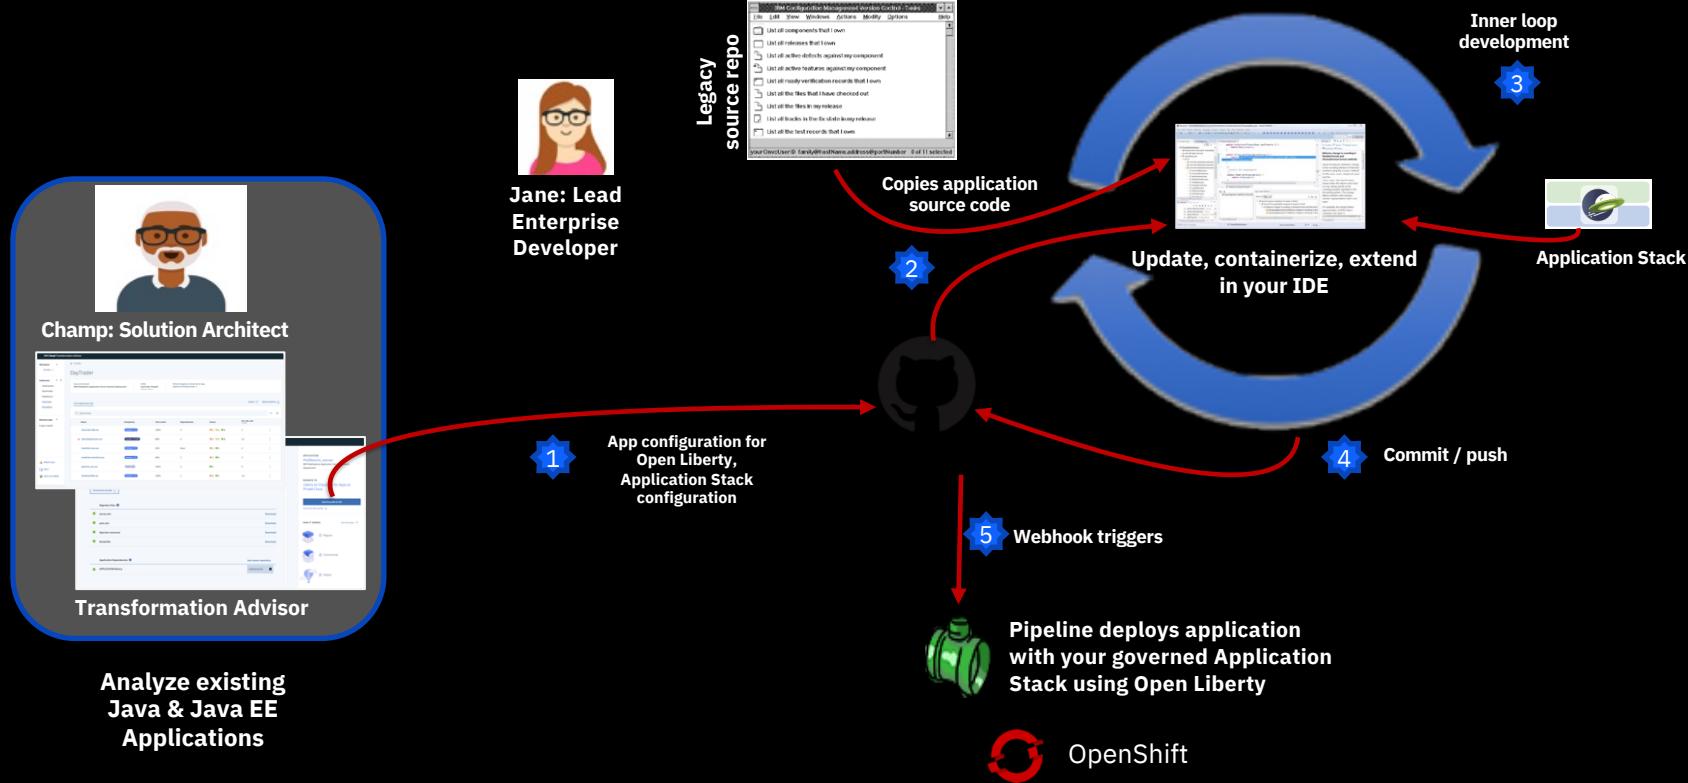
At the bottom, there is a "Send your bundle to Git to begin building and deploying your applications. You can also download the bundle below." section with a "Send to Git" button, which is also highlighted with a blue border. Other buttons in this section include "Download" and "Send to Git".

App Mod developer binary project workflow



App Mod developer source project workflow

*One development experience for both cloud-native and application modernization
lowering the on-ramp to cloud development, test, deployment for legacy app developers*



WebSphere Base Migration Plans

Custom configuration and deployment artifacts created

Migration bundle pushed to git

The screenshot shows the IBM Cloud Transformation Advisor interface. At the top, there's a navigation bar with 'IBM Cloud Transformation Advisor' and a workspace dropdown set to 'Jul8'. Below the navigation is a main content area with tabs for 'Source environment' (selected), 'Destinations', and 'Migration tasks in Cloud Pak for Applications'. A blue box highlights the 'Upload options' button in the top right corner.

The main content area has a title 'WAS' and a sub-section 'Migration bundle'. It displays application details: 'Application name: DefaultApplication.ear.ear', 'Source environment: IBM WebSphere Application Server', and 'Migration target: IBM Cloud Pak for Applications: WebSphere Traditional'. To the right, there's a sidebar titled 'Migration plan' with a tree view showing 'business app' and 'migration plan' nodes, with a blue box highlighting the 'migration plan' node.

Below the application details, there's a section for 'Application dependencies' with two options: 'Manual upload' (selected) and 'Maven repository'. A 'Detected dependencies' table shows one entry: 'Application binary' with a 'Drag or add file' input field. A blue box highlights this table.

At the bottom left, a box highlights the 'Migration Files' section, which lists generated files: 'server_config.py', 'install_app.py', 'Dockerfile', and 'Pipeline resources'. At the bottom right, a box highlights the 'Send to Git' button in the 'How to send files to Git' section, along with 'Download' and 'Send to Git' buttons.

Get WebSphere Application Migration Toolkit

https://developer.ibm.com/wasdev/downloads/#asset/tools-WebSphere_Application_Server_Migration_Tool

Eclipse IDE plugin (Eclipse, WDT, RAD)

Executes source code analysis and provides developer assistance with remediations

WebSphere to Liberty

WebLogic, JBoss, Tomcat to Liberty

WebSphere version-to-version

[+ ALL ASSETS](#)



WebSphere Application Server Migration Toolkit

ASSET TYPE: TOOL

The migration toolkit provides Eclipse-based tools for WebSphere migration scenarios including Cloud migration, WebSphere version to version migration including WAS Liberty, and migration from third-party application servers.

[Download](#)

The migration toolkit provides a rich set of tools that help you migrate applications from third-party application servers, between versions of WebSphere Application Server, to Liberty, and to cloud platforms such as Liberty for Java on IBM Cloud, IBM WebSphere on IBM Cloud and Docker. The following tools are provided:

- Cloud Migration Tool
- WebSphere Version to Version Application Migration Tool
- Apache Tomcat to WebSphere Application Migration Tool
- JBoss to WebSphere Application Migration Tool
- Oracle to WebSphere Application Migration Tool
- WebLogic to WebSphere Application Migration Tool
- Apache Tomcat to Liberty Configuration Migration Tool
- WebSphere Configuration Migration Tool: JBoss
- WebSphere Configuration Migration Tool: WebLogic
- WebSphere Configuration Migration Tool: WebSphere to Liberty

Install it in Eclipse by doing the following:

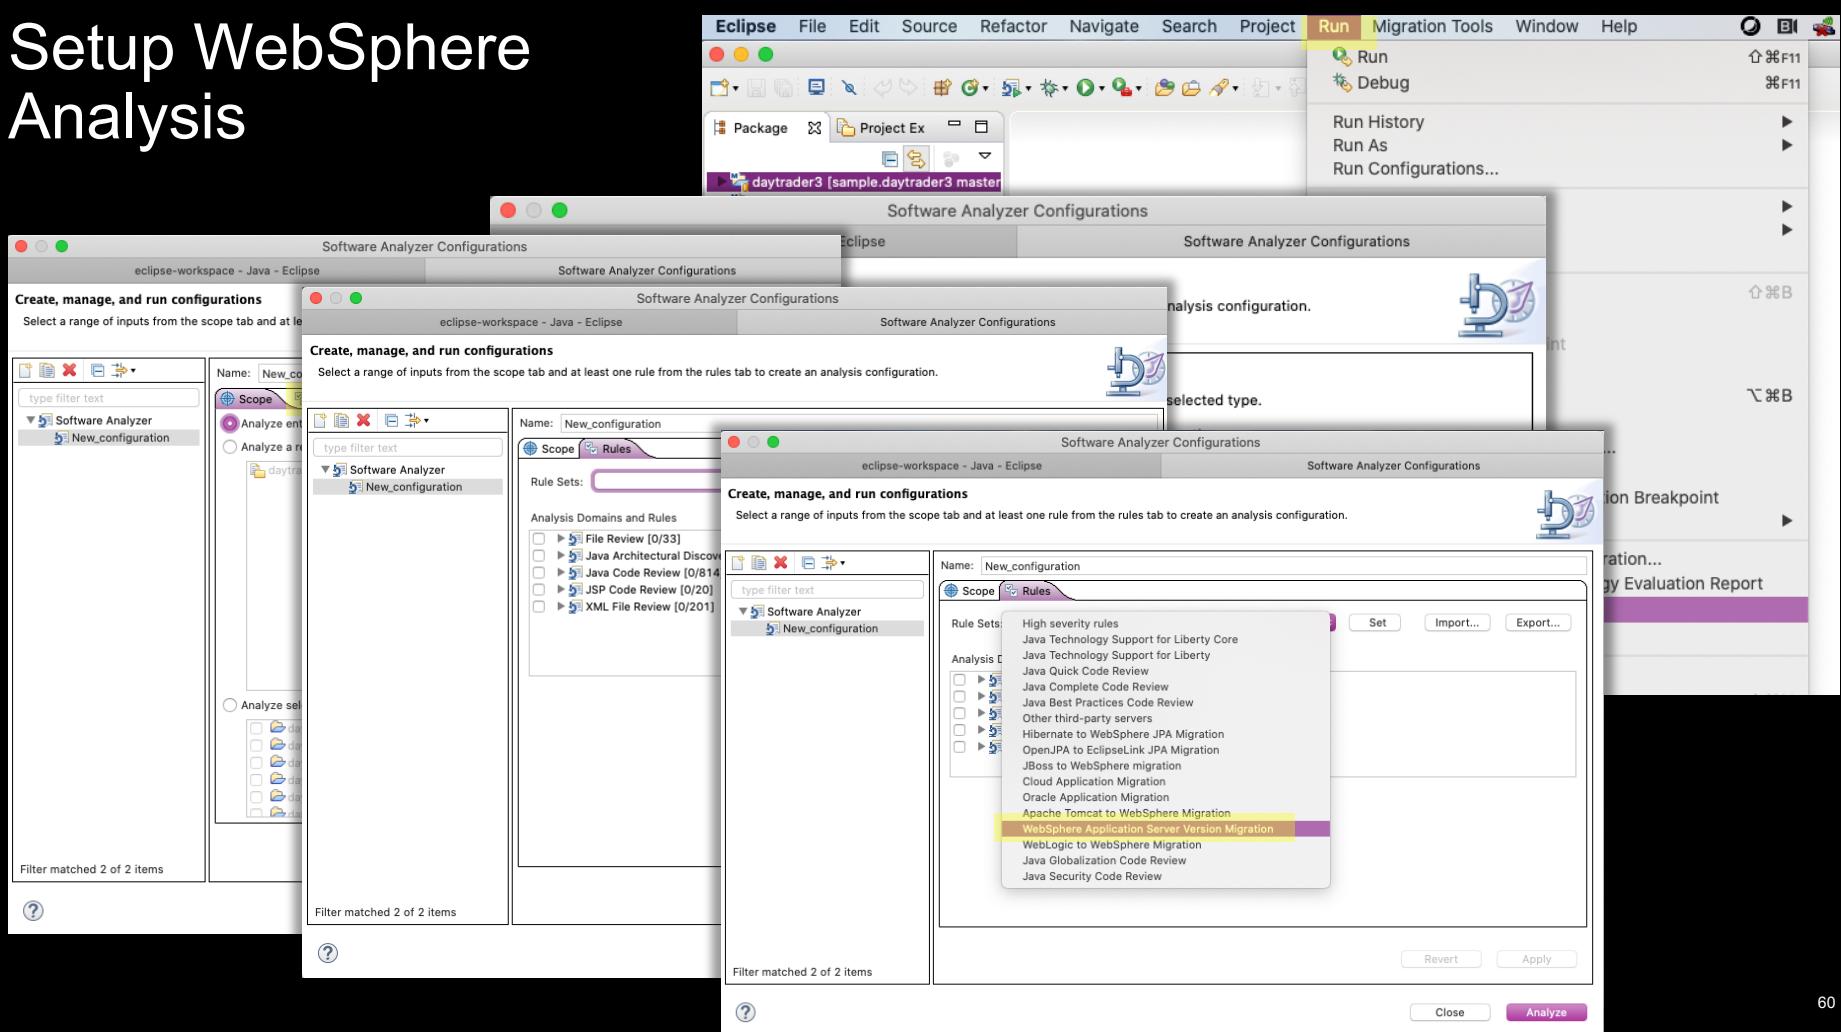
- If you don't already have Eclipse, install [Eclipse IDE for Java EE Developers \(Photon - 4.8 \)](#).

Note: The Business Intelligence, Reporting and Charting (BIRT) dependency is no longer packaged in Eclipse starting with version 4.14 (2019-12). In order to install BIRT on Eclipse 4.14 and later, you will need to download it from the Business Intelligence folder through an older update site such as Eclipse 2019-09 (<http://download.eclipse.org/releases/2019-09>). That site URL will need to be provided in eclipse under **Help > Install New Software**. You will then be able to select and install BIRT.

- To install the migration toolkit, use one of the following methods:
 - Drag this install icon  onto the title bar of Eclipse IDE for Java EE
 - Developers v4.7 or higher:
 - From Eclipse, select **Help > Install New Software**, then click **Add** to add the **repository Location URL**:
`https://public.dhe.ibm.com/ibmdl/export/pub/software/websphere/wasdev/uploads/wamt/MigrationToolkit/`

59

Setup WebSphere Analysis



Setup Source and Target for Analysis

The screenshot shows the WebSphere Application Server Migration Tool interface with several overlapping dialog boxes for rule set configuration.

Main Window: Shows configuration for the "WebSphere Application Server Migration Tool rule selection details". Fields include:

- Source application server: WebSphere Application Server
- Target application server: Liberty Core
- Source Java EE version: Java EE 6
- Target Java EE version: Java EE 8
- Target cloud runtime: None
- Source Java version: IBM Java 6
- Target Java version: IBM Java 8

Top Dialog (Rule set configuration): Shows configuration for "WebSphere Application Server Migration Tool rule selection details". Fields include:

- Source application server: WebSphere Application Server
- Target application server: Liberty
- Source Java EE version: Java EE 6
- Target Java EE version: Java EE 8

Second Dialog (Rule set configuration): Shows configuration for "WebSphere Application Server Migration Tool rule selection details". Fields include:

- Source application server: WebSphere Application Server traditional V8.5.5
- Target application server: Liberty
- Source Java EE version: Java EE 6
- Target Java EE version: Java EE 8

Third Dialog (Rule set configuration): Shows configuration for "WebSphere Application Server Migration Tool rule selection details". Fields include:

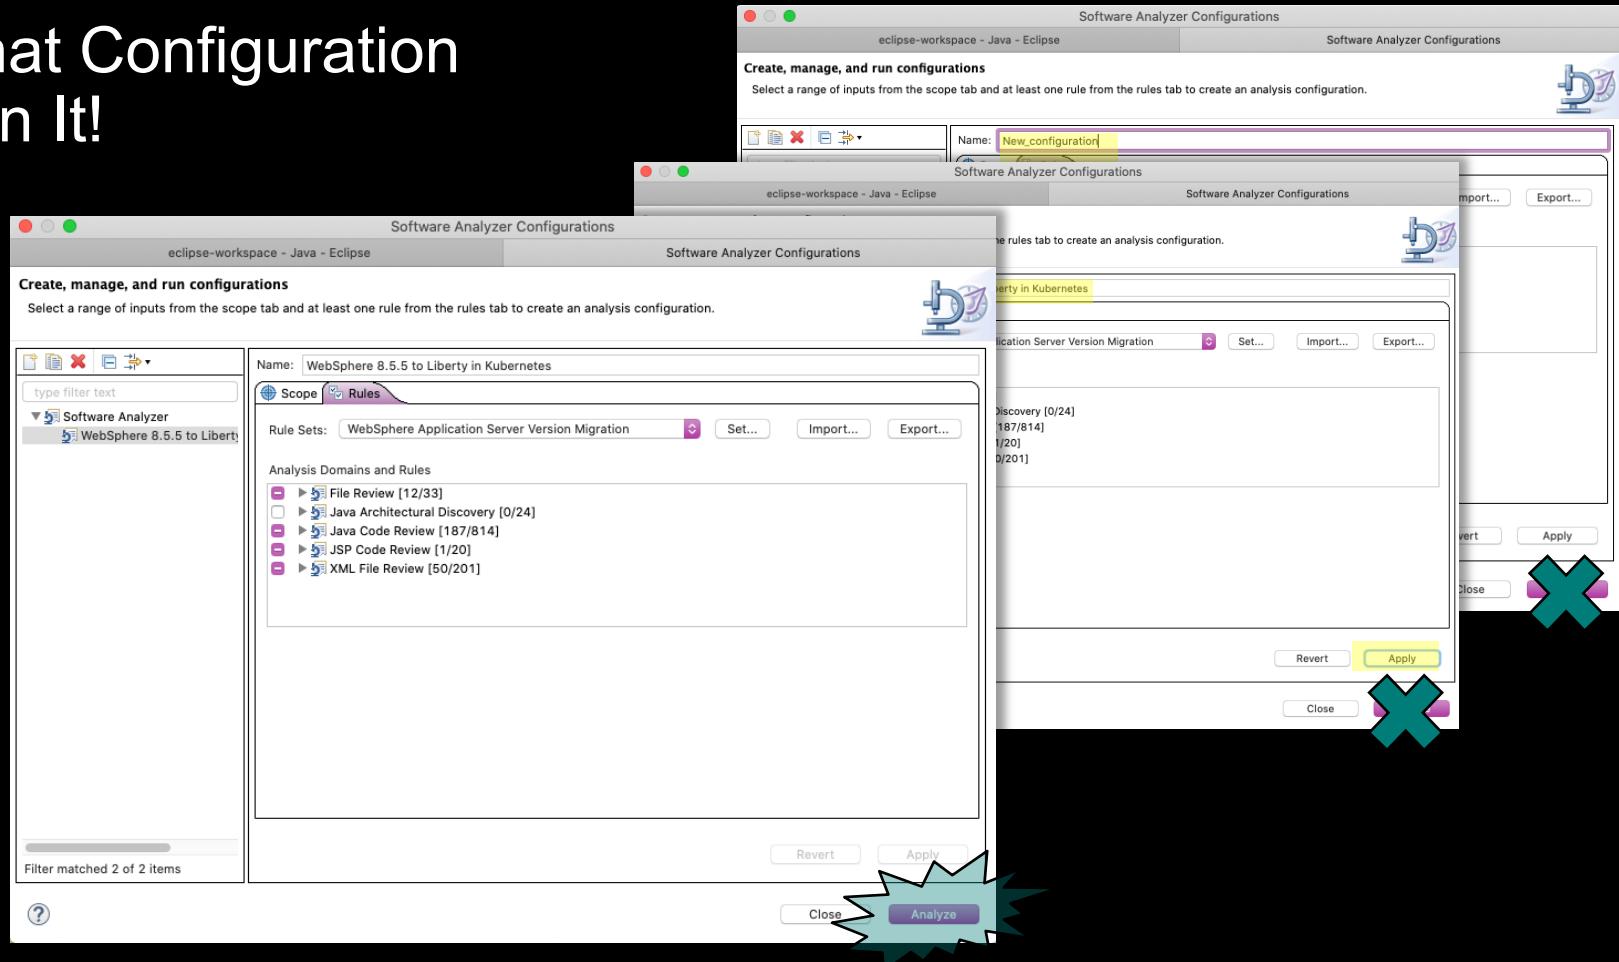
- Source application server: WebSphere Application Server traditional V8.5.5
- Target application server: Liberty
- Source Java EE version: Java EE 6
- Target Java EE version: Java EE 8

Technology Selection: A dropdown menu under "Java EE 8 Technologies" shows the following options:

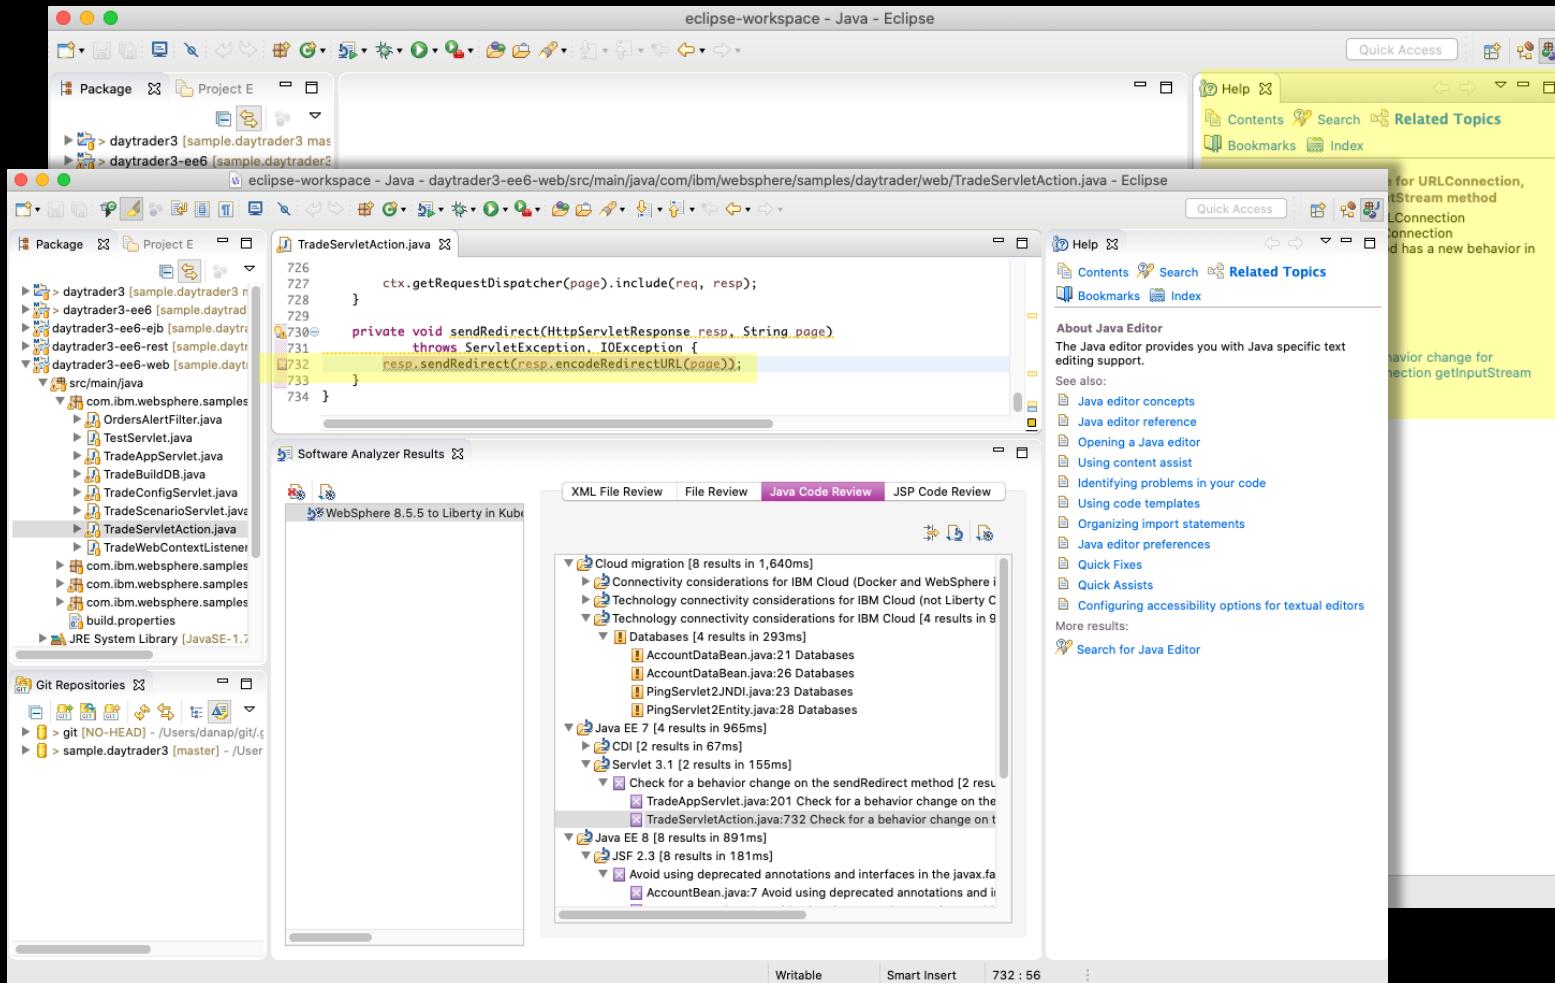
- None
- Docker (IBM Cloud Kubernetes Service) (selected)
- IBM Cloud Runtimes (CF PaaS)
- Third-party PaaS (CF, OSE, etc.)
- WebSphere in IBM Cloud (Virtual Machines)
- IBM Java 8

Buttons: OK, Cancel, Apply, and a large red X button.

Save that Configuration and Run It!



Remedy Issues



Learn More

<https://www.ibm.com/demos/collection/IBM-Cloud-Pak-for-Applications/>



Cloud Pak for Applications

Cloud Pak for Applications is an integrated solution that helps to modernize your applications to cloud native deployments under an offering with flexible entitlement

Cloud Pak for Applications combines application modernization and cloud native application development into a single, consolidated offering. It enables clients to build cloud native applications and to modernize existing applications in a common, more efficient management model. With Cloud Pak for Applications, you get IBM's leading platform solutions in one bundle,

[Product Tours ↓](#)

[Hands-on Labs ↓](#)

Cloud Enabled Use Case:
How to Move an Existing
WebSphere Workload to
Cloud Using IBM WAS VM
Quickstarter

Cloud Enabled Use Case:
Evaluate an on-prem
WebSphere application
for migration

Cloud Enabled Use Case:
How to Deploy Your
Traditional WAS App to
WAS on IBM Cloud Private
using WAS Helm Chart

Cloud Enabled Use Case:
Introduction to Analyzing
Enterprise Applications

Cloud Native Use Case:
Create and Deploy a Cloud
Native App to OpenShift
with Kabanero

Deploy IBM Cloud Pak for
Applications on Managed
OpenShift on IBM Cloud

Hands-on Labs for Cloud Pak for Applications

Cloud Enabled Use Case: App Modernization Journey Part 1 - Evaluation

This is Part 1 of the App Modernization Journey lab series. It will walk you through the process to evaluate the existing Java application using IBM Cloud Transformation Advisor. The lab duration is about 45 minutes.

⌚ 45 minutes

Cloud Native Use Case: Open Liberty Guides

The Open Liberty guides are the best way to learn and explore Open Liberty.

⌚ 10 minutes

DevOps Use Case: Deploy a Cloud Native Application with the IBM Integrated DevOps Pipeline

This lab teaches you how to create a cloud native application and deploy it to a Kubernetes cluster through an integrated DevOps pipeline. The duration of this lab is about 60 minutes.

⌚ 60 minutes

Multi-Cloud Management Use Case: Application Navigator Introduction

This lab walks you through the process to use the IBM Application Navigator tool to manage your cloud native and hybrid applications, across both cloud and traditional platforms. The duration of this lab is about 45 minutes.

⌚ 45 minutes

Microservices application architecture

Cloud native

An application architecture designed to use the strengths and accommodate the challenges of a standardized cloud environment, including:

- Elastic scaling
- Immutable deployment
- Disposable instances
- Less predictable infrastructure

What it means to be cloud native

- Clean contract with underlying OS to ensure maximum portability
- Scale elastically without significant changes to tooling, architecture, or development practices
- Resilient to inevitable failures in the infrastructure and application
- Instrumented to provide both technical and business insight
- Use cloud services such as storage, queuing, and caching
- Rapid and repeatable deployments to maximize agility
- Automated setup to minimize time and cost for new developers

Microservices: Using technology more efficiently

- Microservices is an application architectural style
 - The application is composed of microservice components
- Microservice, a component in this architecture
 - Each is a miniature application
 - Each is focused on one task, a business capability
 - The Single Responsibility principle
 - Each can be deployed and updated independently
 - They are loosely coupled
 - Each has a well-defined interface
 - REST APIs

Microservices: Making developers more efficient

- Each microservice is developed and deployed by a small team
 - The team owns the entire lifecycle of the service
- Microservices accelerate delivery
 - Minimize communication and coordination between people
 - Reducing the scope and risk of change
- Microservices facilitate agile development
 - They make their development teams loosely coupled

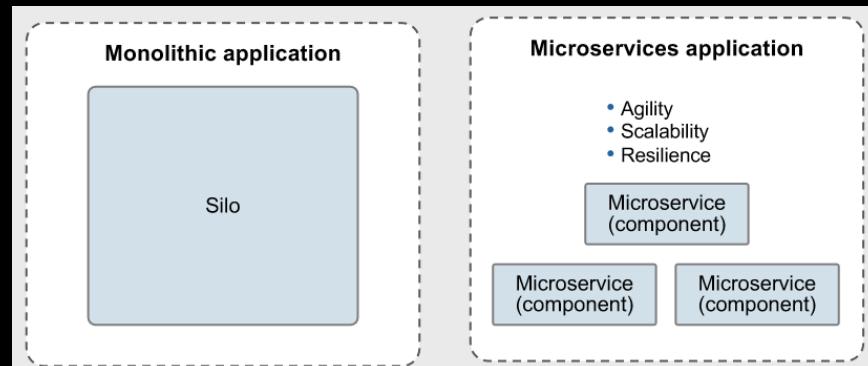
Microservices are an evolution

Evolution of architectural styles

- Monolithic
One large application that does everything
- Microservices
Several smaller applications that each does part of the whole

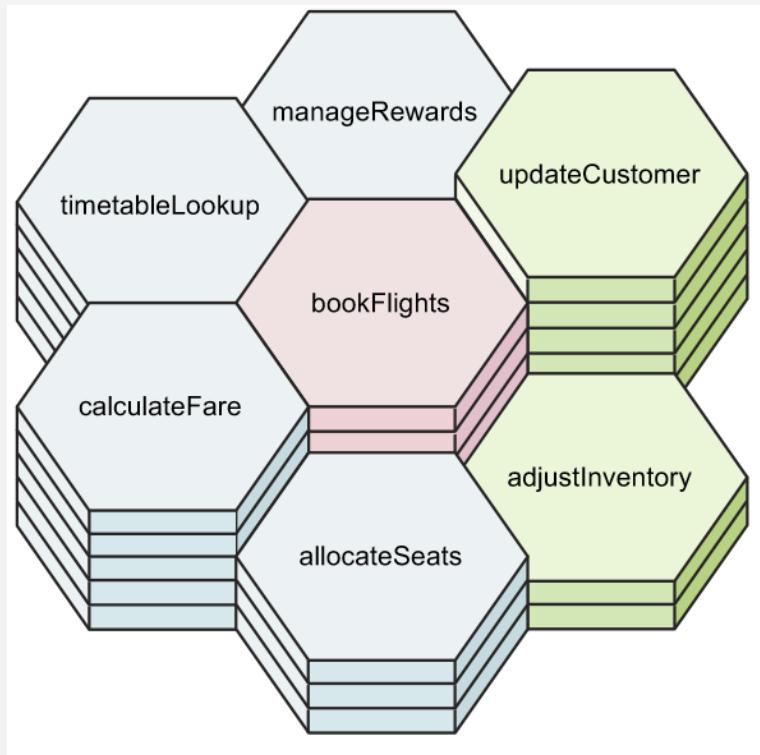
Evolution of service orientation

- SOA
Focused on reuse, technical integration issues, technical APIs
- Microservices
Focused on functional decomposition, business capabilities, business APIs



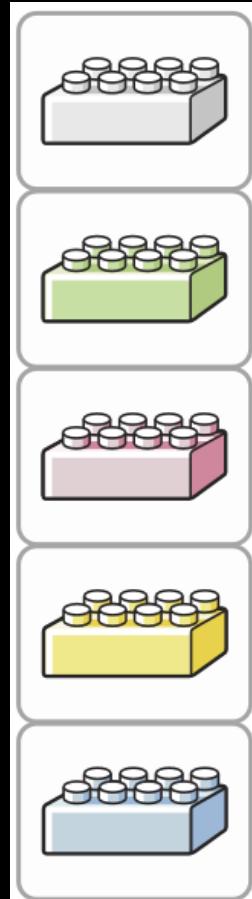
Sample application that uses microservices

- Airline reservation application
 - Seven services (in this example)
- Each service includes
 - Logging
 - Metrics
 - Health check
 - Service endpoint
 - Service registry
 - Service management
- What do the different colors of the tiles mean (red, blue, and green)?
- The tiles are in stacks, some higher than others. What does the height mean?



Key tenets of a microservices architecture

1. Large monoliths are broken down into many small services
 - Each service runs in its own process
 - One service per container
2. Services are optimized for a single function
 - There is only one business purpose per service
 - The Single Responsibility Principle
3. Communication via REST API and message brokers
 - Avoid tight coupling introduced by communication through a database
4. Per-service continuous integration and continuous deployment (CI/CD)
 - Services evolve at different rates
 - You let the system evolve but set architectural principles to guide that evolution
5. Per-service high availability (HA) and clustering decisions
 - One size or scaling policy is not appropriate for all
 - Not all services need to scale; others require autoscaling up to large numbers



Advantages of microservices

- Developed independently
 - Limited, explicit dependencies on other services
- Developed by a single team
 - The team is small
 - All team members can understand the entire code base
- Developed on its own timetable
 - New versions delivered independently of other services
- Polyglot: Each can be developed in a different language
 - Select the best language
- Manages its own data
 - Select the best technology and schema
- Scales and fails independently
 - Isolates problems

Comparing monolithic and microservices architectures

Category	Monolithic architecture	Microservices architecture
Architecture	Built as a single logical executable	Built as a suite of small services
Modularity	Based on language features	Based on business capabilities
Agility	Changes require building and deploying a new version of the entire application	Changes can be applied to each service independently
Scaling	Entire application scaled when only one part is the bottleneck	Each service scaled independently when needed
Implementation	Typically entirely developed in one programming language	Each service can be developed in a different programming language
Maintainability	Large code base is intimidating to new developers	Smaller code bases easier to manage
Deployment	Complex deployments with maintenance windows and scheduled downtimes	Simple deployment as each service can be deployed individually, with minimal downtime

Technology advances have made microservices possible

- Ease and feasibility of distributing components
 - Internet, intranet, or network maturity
 - RESTful API conventions or perceived simplicity, and lightweight messaging
- Ease and simplicity of hosting
 - Lightweight runtimes – Examples: Node.js and WebSphere Liberty
 - Simplified infrastructure
 - OS virtualization (hypervisors), containerization (Docker), infrastructure as a service (cloud infrastructure)
 - Workload virtualization (examples: Cloud Foundry, Kubernetes, OpenWhisk, Swarm)
 - Platform as a service
 - Autoscaling, SLA management, messaging, caching, build management
- Agile development methods
 - Examples: IBM Cloud Garage Method, XP, TDD, Scrum, Continuous Delivery
 - Standardized code management, such as GitHub

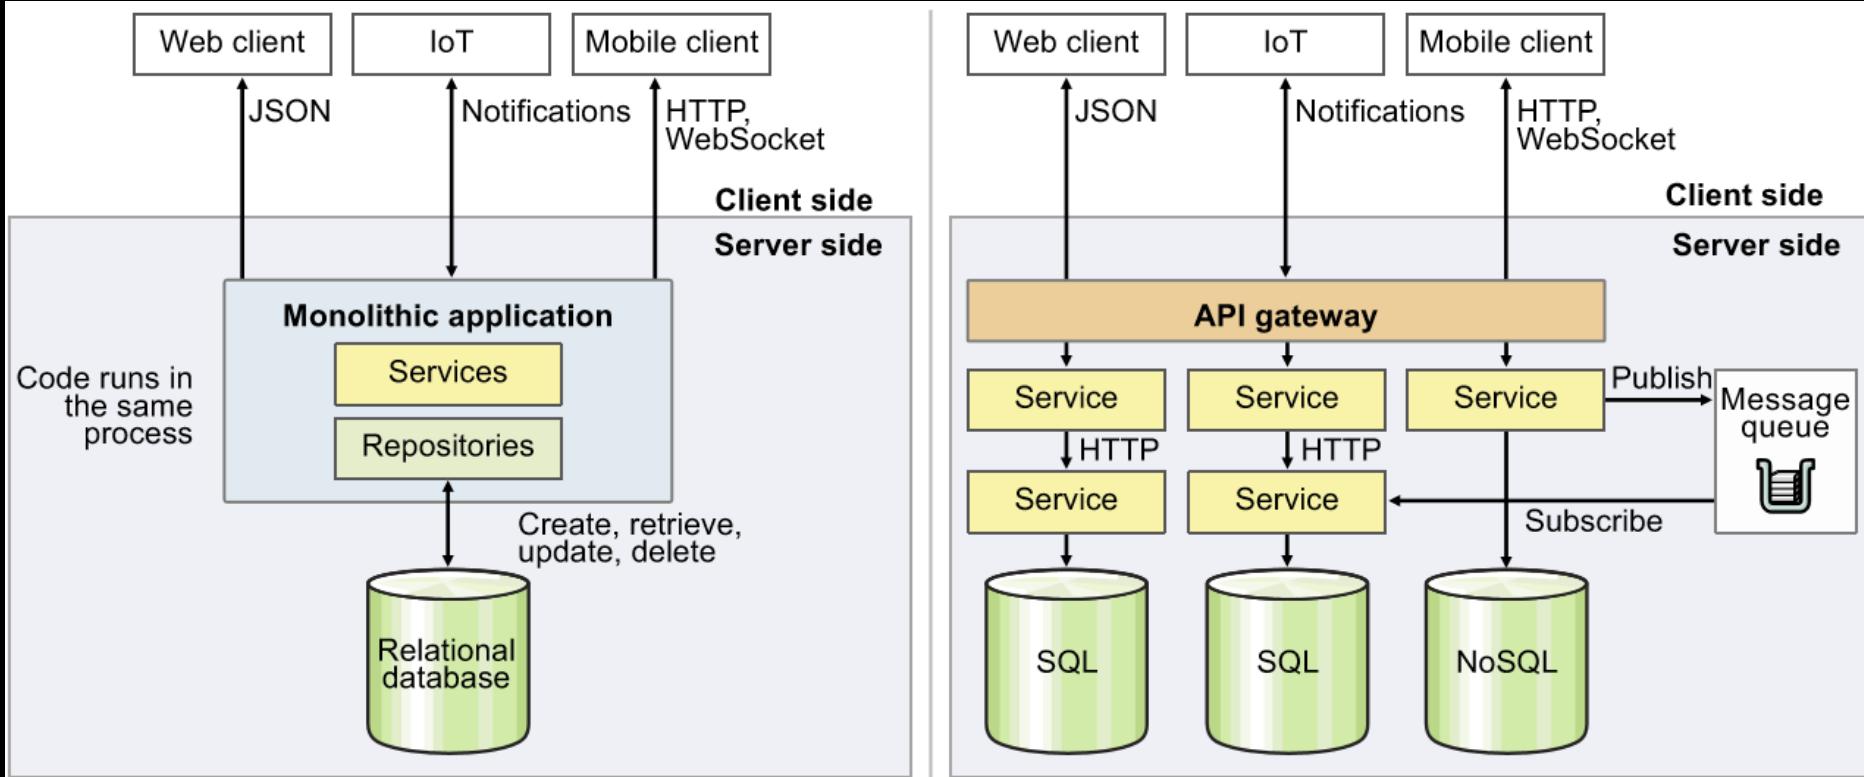
Microservice challenges

- Greater operational complexity because there are more moving parts to monitor and manage
- Developers must have significant operational skills (DevOps)
- Service interfaces and versions
- Duplication of effort across service implementations
- Extra complexity of creating a distributed system
 - Network latency
 - Fault tolerance
 - Serialization
- Designing decoupled, non-transactional systems is difficult
- Locating service instances
- Maintaining availability and consistency with partitioned data
- End-to-end testing

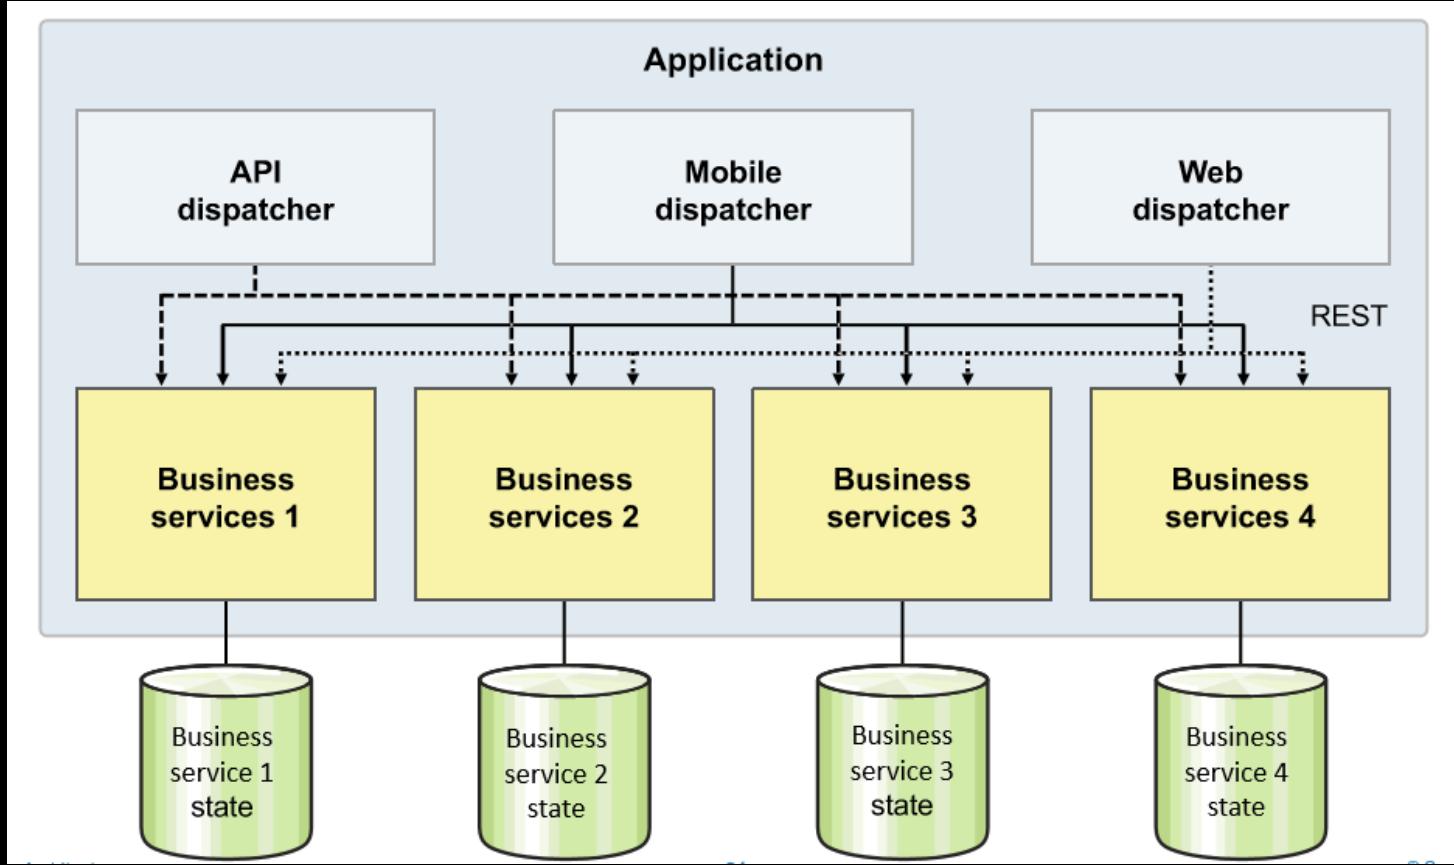
Microservices and SOA

- Both SOA and microservices deal with a system of services that communicate over a network But there are differences
- The focus of SOA is on reuse
 - This tends to align with a centrally funded model
 - SOA services tend to be “servants of many masters”
 - This means that a change to a SOA service might impact multiple consumers
- The focus of microservices is on breaking down a potentially monolithic application into smaller, more manageable components
 - With the objective of more flexible, decoupled, faster development
 - Challenges here relate to needing good DevOps, management views, and controls

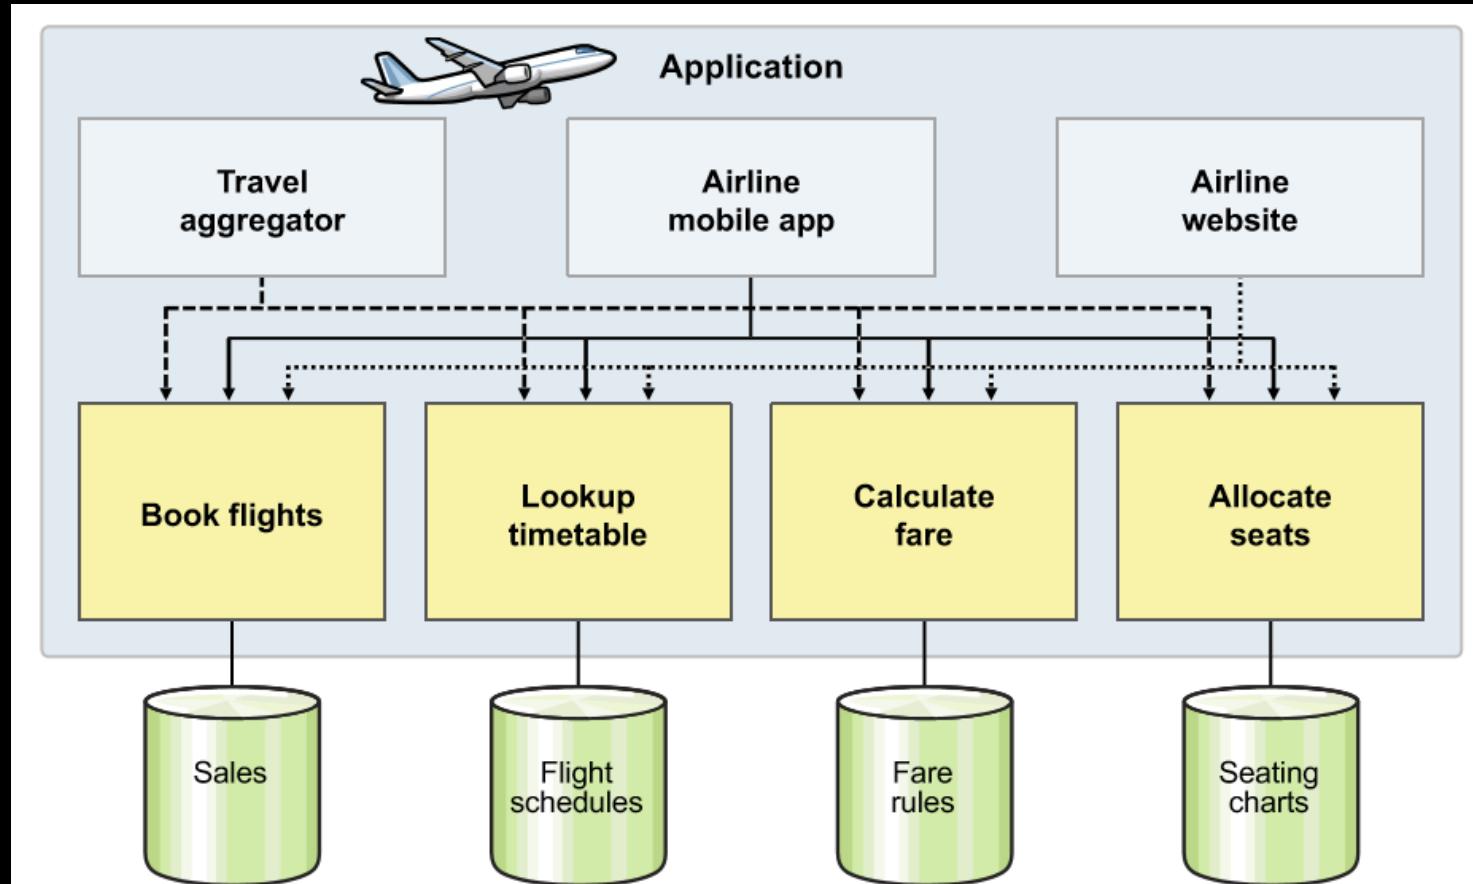
Monolithic architecture versus microservices architecture



Microservices architecture

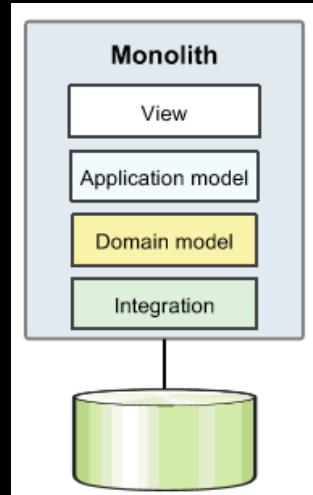
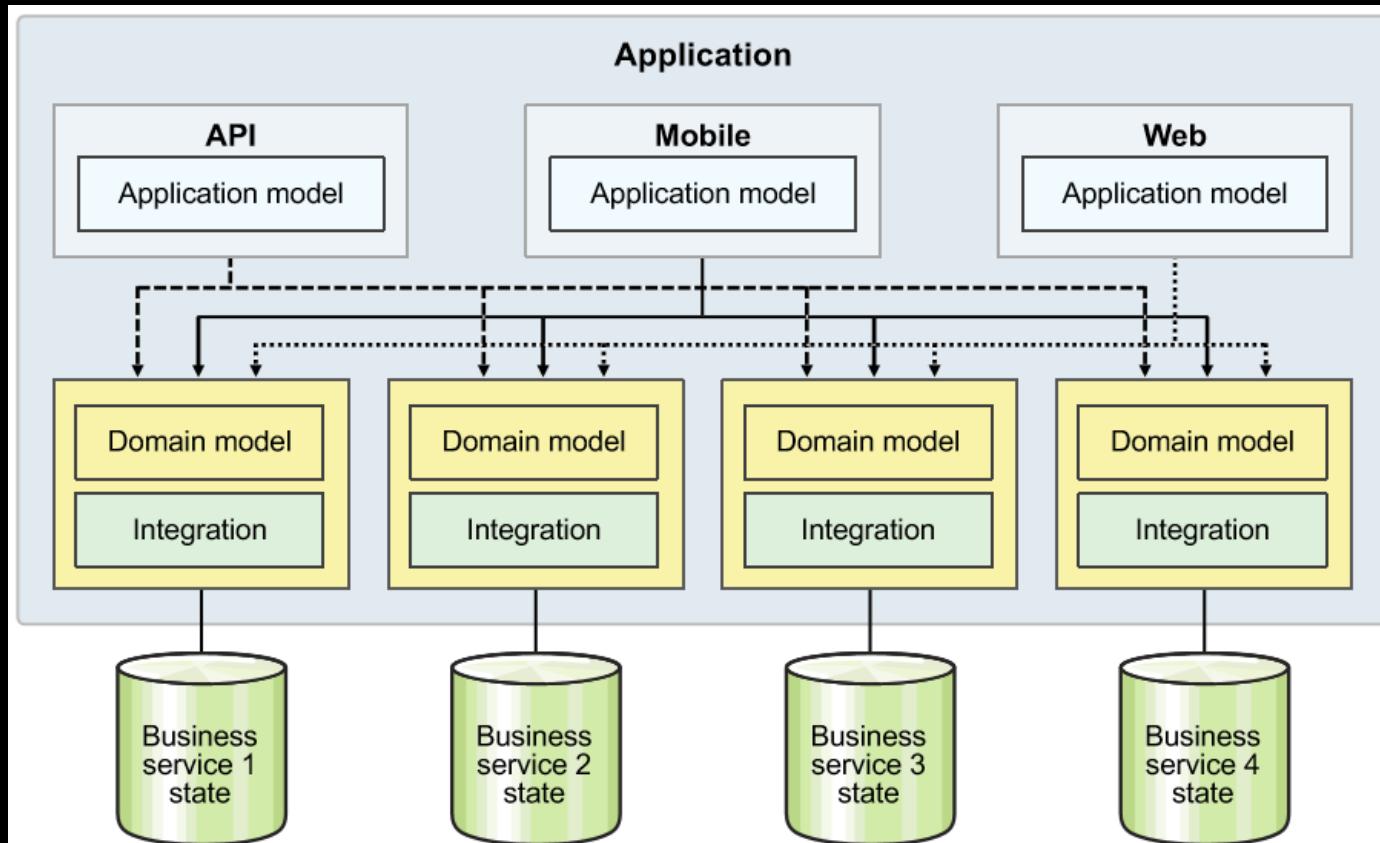


Sample airline architecture



Microservices layers

Each service is a smaller half-monolith

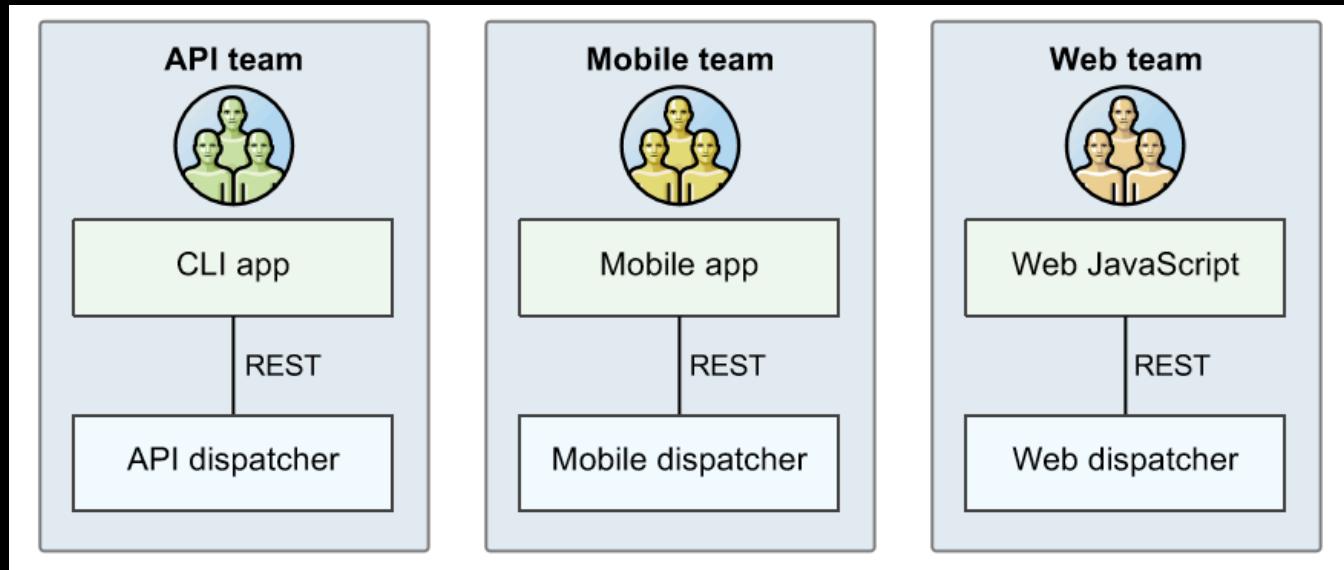


Language decisions

- Dispatchers are most often written in Node.js
 - Better fidelity with the clients
 - Especially clients that run JavaScript
 - iOS teams might want to use the Swift server-side runtime
 - I/O intensive, supports numerous concurrent clients
- Business services are most often written in Java
 - Better for CPU-intensive tasks
 - Better connectivity to external systems
- The team selects the language
 - Choose the language that best fits the job

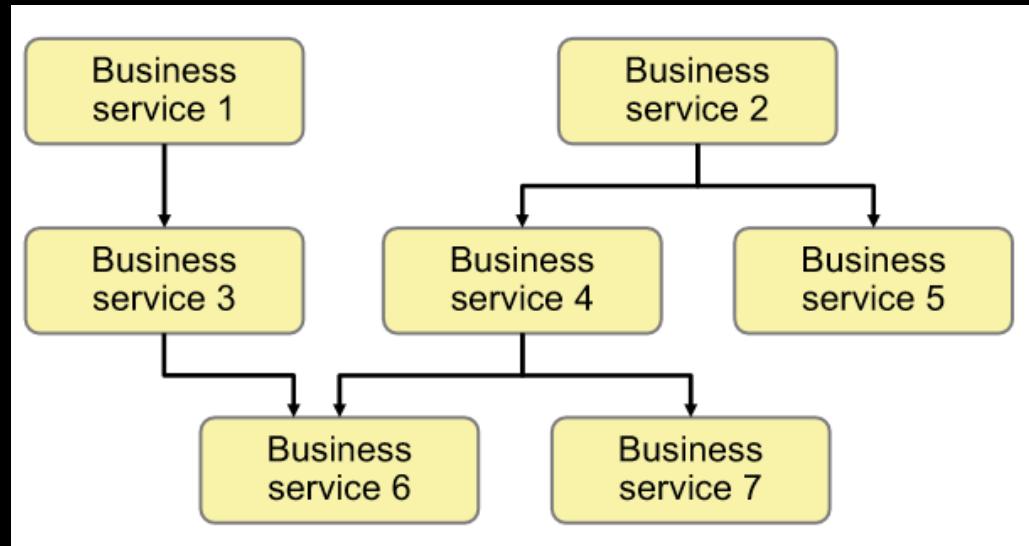
Back ends for front ends

- Each dispatcher is a back end for an external front end, typically a GUI
- The same team develops each back end and front end pair
 - Use the same or compatible languages in the pair



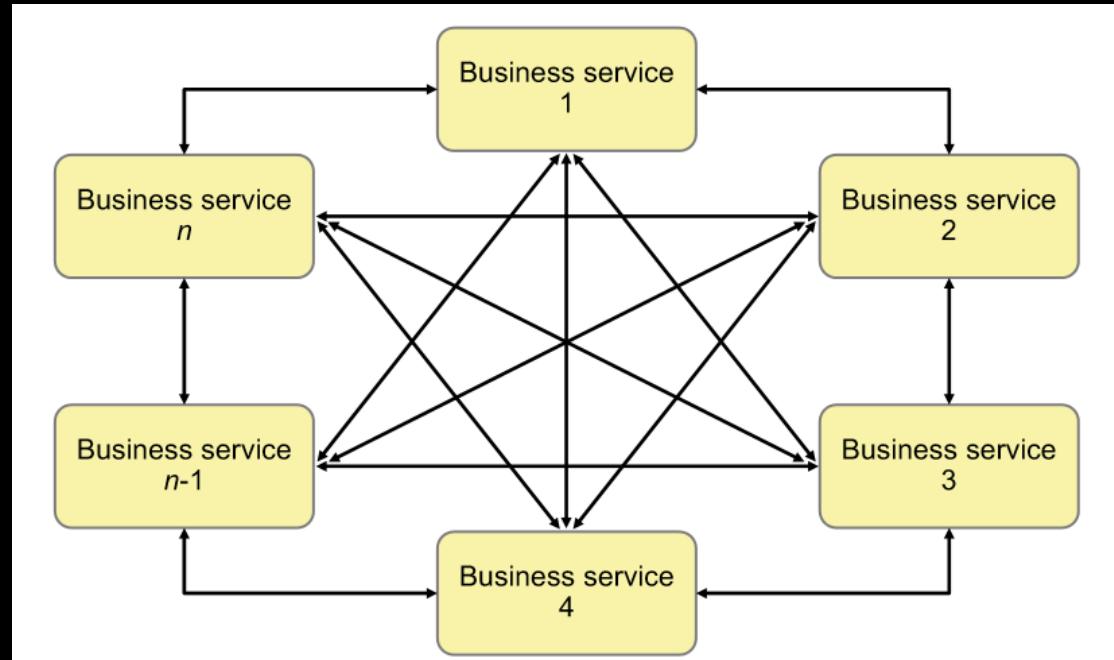
Business service microservices dependencies: Typical

- Business services can delegate to other business services
 - Avoid circular dependencies
- Be careful that each service still implements a complete task
- They are not separate layers



Business service microservices dependencies: Death Star

- Practically every service delegates to or coordinates with every other service
- Difficult to deploy and maintain



Twelve factor applications

The twelve-factor app

- A set of best practices for creating applications
 - Implementing, deploying, monitoring, and managing
- Typical modern applications
 - Deployed in the cloud
 - Accessible as web applications that deliver software-as-a-service (SaaS)
- Can be applied to any application
 - Implemented in any programming language
 - Using any backing services such as database, messaging, and caching
- Addresses common problems
 - The dynamics of the growth of an app over time
 - The dynamics of collaboration between developers
 - Avoiding the cost of software erosion
 - Systemic problems in modern application development
- Provides a shared vocabulary for addressing these problems

The twelve factors

- I. **Codebase**: One codebase that is tracked in revision control, with many deployments
- II. **Dependencies**: Explicitly declare and isolate dependencies
- III. **Configuration**: Store Configuration in the environment
- IV. **Backing services**: Treat backing services as attached resources
- V. **Build, release, run**: Strictly separate build and run stages
- VI. **Processes**: Execute the app as one or more stateless processes
- VII. **Port binding**: Export services with port binding
- VIII. **Concurrency**: Scale out using the process model
- IX. **Disposability**: Maximize robustness with fast startup and efficient shutdown
- X. **Development and production parity**: Keep development, staging, and production as similar as possible
- XI. **Logs**: Treat logs as event streams
- XII. **Admin processes**: Run administrative and management tasks as one-off processes

Factor 1: Codebase

I. Codebase

- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- One codebase tracked in source code management (SCM) with versioning
- Multiple deployments from the same codebase

Factor 2: Dependencies

- I. Codebase
- II. Dependencies**
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Explicitly declare and isolate dependencies
- Typically language-specific
 - Node.js: Node Package Manager (NPM)
 - Liberty: Feature manager
 - Ruby: Bundler
 - Java EE: Application resources
- Never rely on system-wide dependencies

Factor 2: Dependencies

- I. Codebase
- II. Dependencies**
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Explicitly declare and isolate dependencies
- Typically language-specific
 - Node.js: Node Package Manager (NPM)
 - Liberty: Feature manager
 - Ruby: Bundler
 - Java EE: Application resources
- Never rely on system-wide dependencies

Dependencies: Liberty feature manager

- Server configuration is part of deploying the application
- You manage it in source control

```
<server description="Portfolio server">
  <featureManager>
    <feature>microProfile-3.2</feature>
    <feature>jdbc-4.2</feature>
    <feature>jms-2.0</feature>
    <feature>jca-1.7</feature>
    <feature>jndi-1.0</feature>
    <feature>appSecurity-2.0</feature>
    <feature>monitor-1.0</feature>
    <feature>jpa-2.2</feature>
    <feature>localConnector-1.0</feature>
    <feature>mpReactiveMessaging-1.0</feature>
    <feature>ejbLite-3.2</feature>
  </featureManager>
```

Factor 3: Configuration

- I. Codebase
- II. Dependencies
- III. Configuration**
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Store configuration in the environment
- Separate configuration from source
- Enables the same code to be deployed to different environments

Externalizing Configuration

- Liberty
 - Variables in the configuration that can be taken from properties or environment entries

```
<dataSource id="PortfolioDB" jndiName="jdbc/Portfolio/PortfolioDB">
    <jdbcDriver>
        <library description="DB2 JDBC driver jar" name="DB2">
            <file id="db2jcc4" name="db2jcc4.jar"/>
        </library>
    </jdbcDriver>
    <properties.db2.jcc databaseName="${JDBC_DB}" portNumber="${JDBC_PORT}" serverName="${JDBC_HOST}"
        user="${JDBC_ID}" password="${JDBC_PASSWORD}" />
</dataSource>
```

- Kubernetes/OpenShift
 - Uses ConfigMap and Secret resources

```
kubectl create secret generic apikey --from-literal=API_KEY=123-456
```

```
kubectl create configmap language --from-literal=LANGUAGE=English
```

Factor 4: Backing services

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services**
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Treat backing services as attached resources:
 - Databases
 - Messaging systems
 - LDAP servers
 - Others
- Local and remote resources should be treated identically
 - Possible locations for run time and resources:
 - In the same process
 - On the same host
 - On different hosts in the same data center
 - In different data centers

Factor 5: Build, release, run

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run**
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

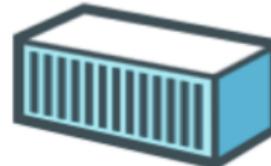
- Strictly separate build and run stages

Docker: Build, ship, run

- Build: Create the runtime that goes inside the container
 - Application
 - Application server
 - Operating system
- Ship: Create a snapshot of the runtime
 - A container image
 - Stored in a registry
- Run: Create a running application
 - A container running in a container engine



Build



Ship



Run

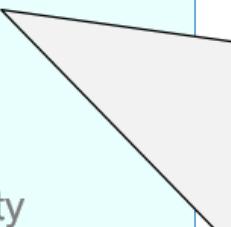
Factor 6: Processes

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes**
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Run the application as one or more stateless processes
- Do not rely on session affinity, also called sticky sessions

Factor 7: Port binding

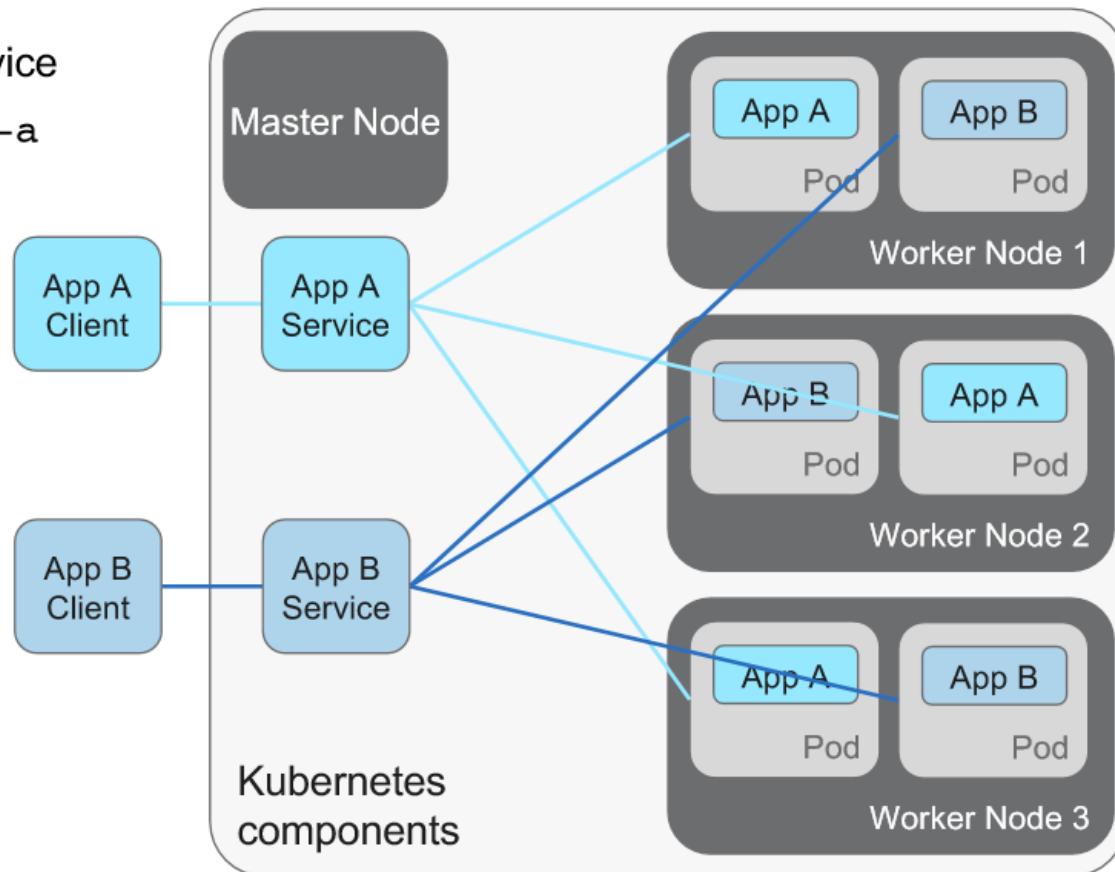
- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding**
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- 
- Export services with port binding
 - Web app binds to an HTTP port and listens for requests coming in on that port

Port binding: Kubernetes service

- Expose a set of pod replicas as a service

```
kubectl expose deployment/app-a  
  --type=LoadBalancer  
  --port=8080  
  --name=app-a-service  
  --target-port=8080
```



Factor 8: Concurrency

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency**
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Scale out using the process model
To add capacity, run more instances
- There are limits to how far an individual process can scale
- Stateless applications make scaling simple

Factor 9: Disposability

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability**
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes

- Maximize robustness with fast startup and efficient shutdown
- Application instances are disposable
- Application should handle shutdown signal or hardware failure with crash-only design

Factor 10: Development and production parity

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. **Dev/prod parity**
- XI. Logs
- XII. Admin processes

- Keep development, staging, and production as similar as possible
- Use the same backing services in each environment

Factor 11: Logs

I. Codebase

II. Dependencies

III. Configuration

IV. Backing services

V. Build, release, run

VI. Processes

VII. Port binding

VIII. Concurrency

IX. Disposability

X. Dev/prod parity

XI. Logs

XII. Admin processes

- Treat logs as event streams
- Each process writes to stdout
 - Application should not write to specialized log files
 - Environment decides how to gather, aggregate, and persist stdout output

Factor 12: Administrative processes

- I. Codebase
- II. Dependencies
- III. Configuration
- IV. Backing services
- V. Build, release, run
- VI. Processes
- VII. Port binding
- VIII. Concurrency
- IX. Disposability
- X. Dev/prod parity
- XI. Logs
- XII. Admin processes**

Run administrative and management tasks as single processes, such as these examples:

- Tasks for performing database migrations
- Debugging

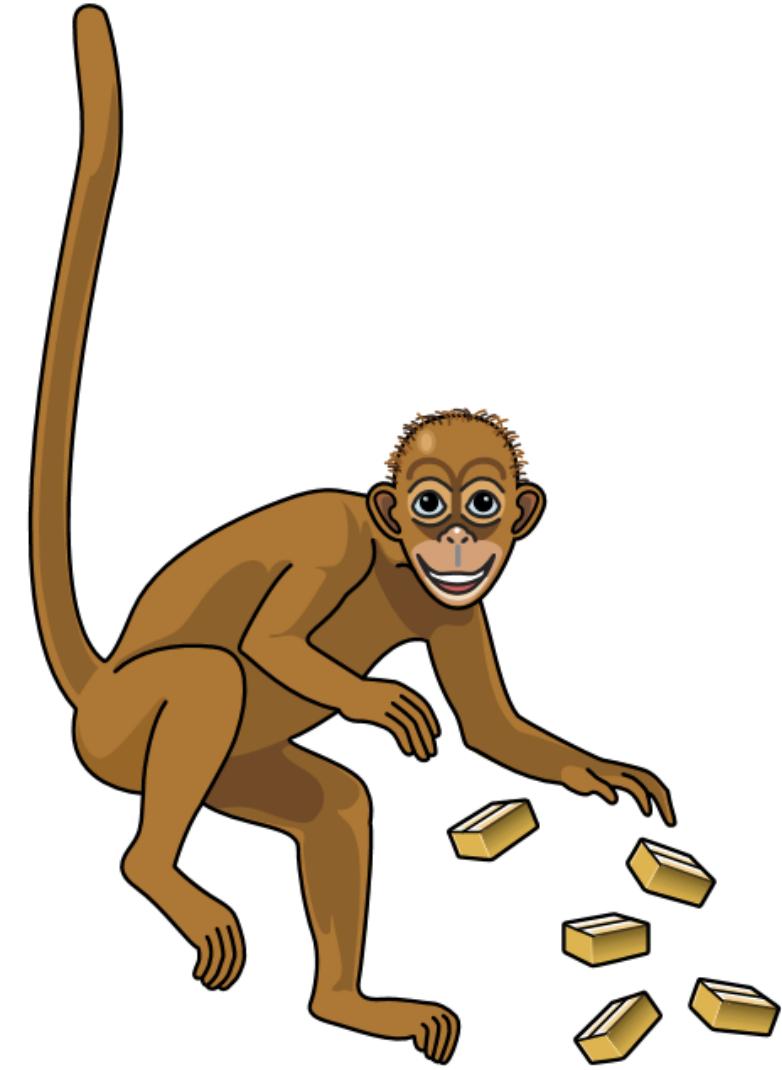
Developing Microservices

Developing a microservices application

- Identify a set of independent business tasks
 - Build initial microservices around those tasks
 - Teams that can work independently
 - Components that can be deployed, scale, and fail independently
- Design for failure
 - Use microservices to make the application more robust
- Design for scale
 - Service discovery
 - For example, Eureka
 - Configuration repositories
 - For example Zookeeper
 - Common logging
- A service mesh is very helpful for steps 2 and 3

Design for failure

- Any service can fail
 - In a monolith, when one part stops working, it all stops working
 - Application must keep working and stay responsive
- Employ patterns for resiliency
 - Service Registry: Dynamic listing of service instances
 - Circuit Breaker: Block calls to a service that's not working
 - Bulkhead: Separate connection pools for separate resources
 - Command: Make requests easy to retry, throttle, and monitor
- Test for failure
 - Purposely, randomly cause problems
 - Example: Simian Army framework



MicroProfile: Java Microservices Programming Model

MicroProfile programming model

- Exposes REST APIs
 - Implemented using JAX-RS
- Uses Contexts and Dependency Injection (CDI)

```
@Path("props")
public class SystemProperties {

    @Produces(MediaType.APPLICATION_JSON)
    @GET
    public Map getProps() {
        return System.getProperties();
    }
}
```

```
@Path("props")
@RequestScoped
public class ServiceC {
    @Inject
    private MyBean bean;
```

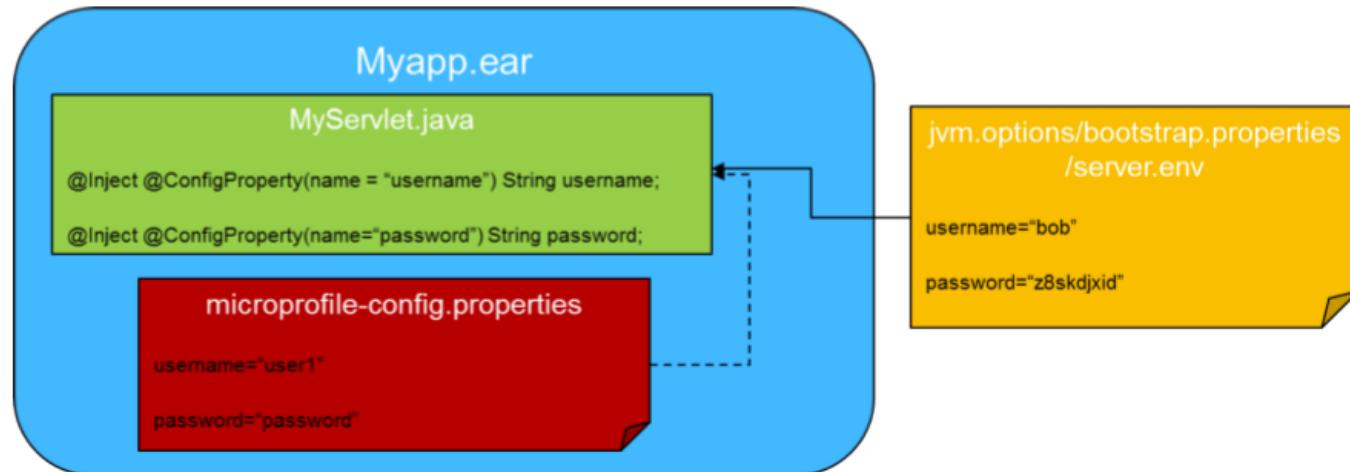
MicroProfile Config feature

Config	Fault Tolerance	Health Check	Health Metrics	JWT Propagation
externalize configuration to improve portability mpConfig-1.0	build robust behavior to cope with unexpected failures	common format to determine service availability	common REST endpoints for monitoring service health	interoperable authentication and role-based access control

- Write once
- Containerize once
- Liberty injects config from outside the container
- Portable across environments

The property file, *microprofile-config.properties*, packaged in the application can be overridden by

1. System variables (400 as the default priority)
2. Environment variables (300 as the default priority) or
3. A custom property files with a higher priority than *microprofile-config.properties* (100 as the default priority)



MicroProfile: Fault Tolerance feature

Config	Fault Tolerance	Health Check	Health Metrics	JWT Propagation
externalize configuration to improve portability	build robust behavior to cope with unexpected failures	common format to determine service availability	common REST endpoints for monitoring service health	interoperable authentication and role-based access control

mpFaultTolerance-1.0

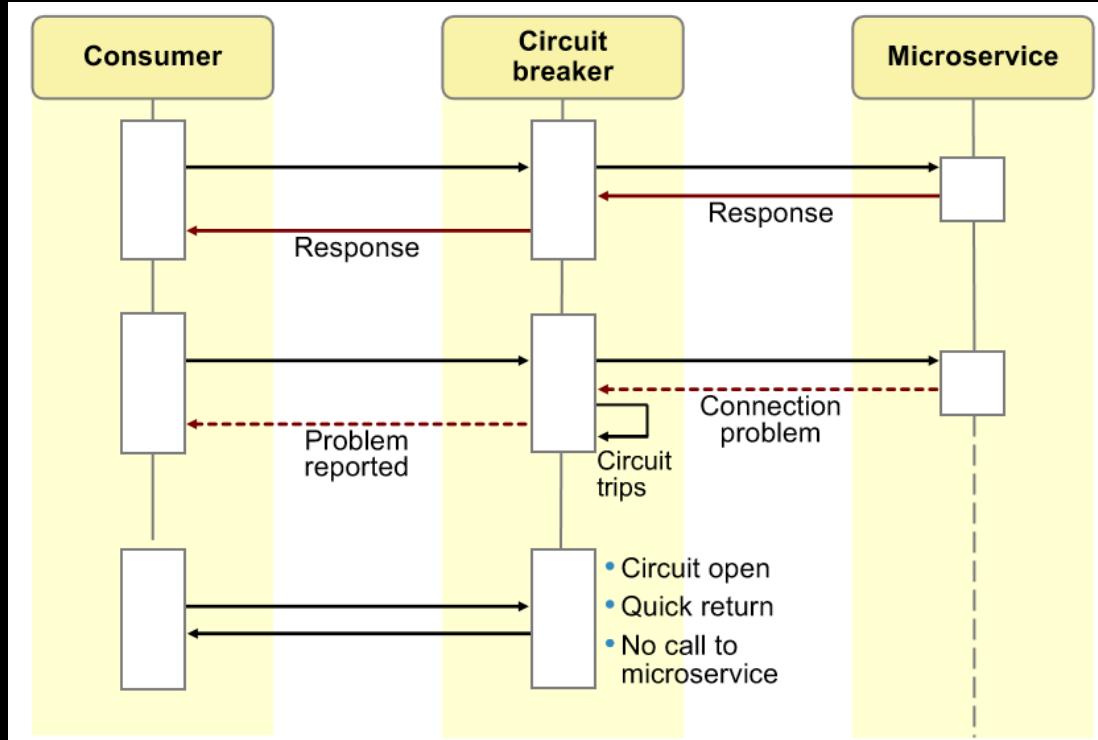
- Declarative policies to make microservices more resilient when runtime failures occur
 - Implemented as Java annotations
 - Support for common application resiliency patterns
 - Timeout, Retry, Circuit Breaker, Bulkhead, Fallback
 - Policies configured by properties
- Defined on the method that invokes the remote microservice
 - Handled by Contexts and Dependency Injection (CDI) interceptors

Fault Tolerance – Annotations

- Asynchronous: Invoke the method asynchronously
 - `@Aynchronous`
- Retry: Invoke a method again after failure
 - `@Retry(maxRetries=3, delay=400, maxDuration=2000)`
- Timeout :Specify a timeout for the method invocation
 - `@Timeout(300)`
- Bulkhead: Limit the concurrent requests
 - `@Bulkhead(20)` – semaphore model, at most 20 concurrent requests
 - `@Asynchronous @Bulkhead(value=20, waitingTaskQueue=30)` - thread pool style, at most 20 concurrent threads
- CircuitBreaker: Open a circuit so that the following invocation will fail immediately
 - `@CircuitBreaker(delay=200, failureRatio=0.2)`
- Fallback: A secondary service when the primary service invocation fails
 - `@Fallback(MyFallbackHandler.class)` - invoke `MyFallbackHandler.handle()` on exception
 - `@Fallback(fallbackMethod="fbMethod")` - invoke `fbMethod()` on the same instance as the method with this annotation

Circuit breaker

- A service consumer is only as reliable as its service provider
- Avoid invoking an unreliable service instance
 - Service might throw an error
 - Service might time out
 - Network might fail
- Circuit breaker
 - Consumer can invoke reliably
 - Invokes unreliable service
 - Fails fast when the service isn't working



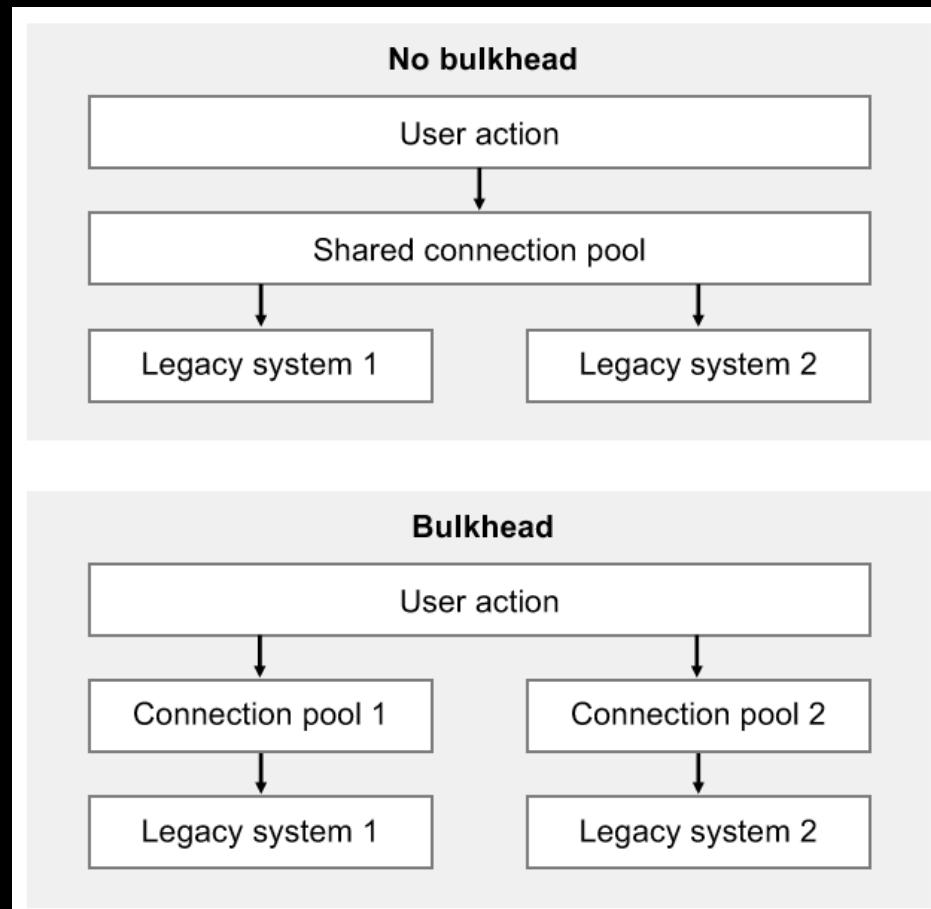
Bulkhead

No bulkhead

- Single connection pool shared to connect to multiple external systems
- If one system blocks connections
 - All connections block on the one system
 - No way to connect to the working systems

Bulkhead

- Separate connection pool for each external system
- If one system blocks connections
 - All its connections block
 - Meanwhile, connections to the working systems still work



MicroProfile: Health Check feature

Config	Fault Tolerance	Health Check	Health Metrics	JWT Propagation
externalize configuration to improve portability	build robust behavior to cope with unexpected failures	common format to determine service availability mpHealth-1.0	common REST endpoints for monitoring service health	interoperable authentication and role-based access control

- Exposes **/health** default endpoint for the server/container if feature enabled
- Offers health checks for both readiness and liveness.
 - Readiness tells whether a microservice is ready to process requests. For example, a readiness check might check dependencies, such as database connections.
 - Liveness check tells whether a microservice is running. If fails, the application can be terminated.

MicroProfile: Health Check Readiness

- The `@Readiness` annotation indicates that this particular bean is a readiness health check procedure. By pairing `@ApplicationScoped`, the bean is discovered automatically when the `/health` endpoint receives a request.
- Access the `/health/ready` endpoint to view the data from the readiness health checks

```
@Readiness
@ApplicationScoped
public class SystemReadinessCheck implements HealthCheck {

    @Override
    public HealthCheckResponse call() {
        ...
        if (notReady) {
            return HealthCheckResponse.down(readinessCheck);
        }
        return HealthCheckResponse.up(readinessCheck);
    }
}
```

MicroProfile: Health Check Liveness

- The `@Liveness` annotation indicates that this particular bean is a liveness health check procedure. By pairing `@ApplicationScoped`, the bean is discovered automatically when the `/health` endpoint receives a request.
- Access the `/health/live` endpoint to view the data from the liveness health checks

```
@Liveness
@ApplicationScoped
public class SystemLivenessCheck implements HealthCheck {

    @Override
    public HealthCheckResponse call() {
        MemoryMXBean memBean = ManagementFactory.getMemoryMXBean();
        long memUsed = memBean.getHeapMemoryUsage().getUsed();
        long memMax = memBean.getHeapMemoryUsage().getMax();

        return HealthCheckResponse.named(SystemResource.class.getSimpleName() + " Liveness Check")
            .withData("memory used", memUsed)
            .withData("memory max", memMax)
            .state(memUsed < memMax * 0.9).build();
    }
}
```

MicroProfile: Health Metrics feature

Config	Fault Tolerance	Health Check	Health Metrics	JWT Propagation
externalize configuration to improve portability	build robust behavior to cope with unexpected failures	common format to determine service availability	common REST endpoints for monitoring service health mpMetrics-1.0	interoperable authentication and role-based access control

- Exposes /metrics endpoint for the server/container if feature enabled
- Exposes system, vendor, and app-specific metrics
- Response in JSON
 - For collection from collectd or other JSON-friendly tools
 - Prometheus text formats
- App metrics provided
 - Dropwizard-based API
 - New CDI-enabled annotations
- Out-of-the-box metrics
 - JVM memory
 - Garbage Collection
 - JVM uptime
 - Threads
 - Thread Pools (stretch goal)
 - ClassLoading
 - CPU usage and availability

MicroProfile: Health Metrics annotations

- `@Counted` – counts invocations of annotated object
- `@Gauge` – samples the value of annotated object
- `@ConcurrentGauge` – counts parallel invocations
- `@Metered` – tracks frequency of invocations
- `@Metric` – contains the metadata information
- `@Timed` – aggregates timing durations and provides duration statistics, plus throughput statistics
- `@SimplyTimed` - only tracks elapsed time duration and count

MicroProfile: Metrics example

```
@Timed(name = "inventoryProcessingTime",
       tags = {"method=list"},
       absolute = true,
       description = "Time needed to process the inventory")
@Counted(name = "inventoryAccessCount",
          absolute = true,
          description = "Number of times the list of systems method is requested")
public InventoryList list() {
    return new InventoryList(systems);
}

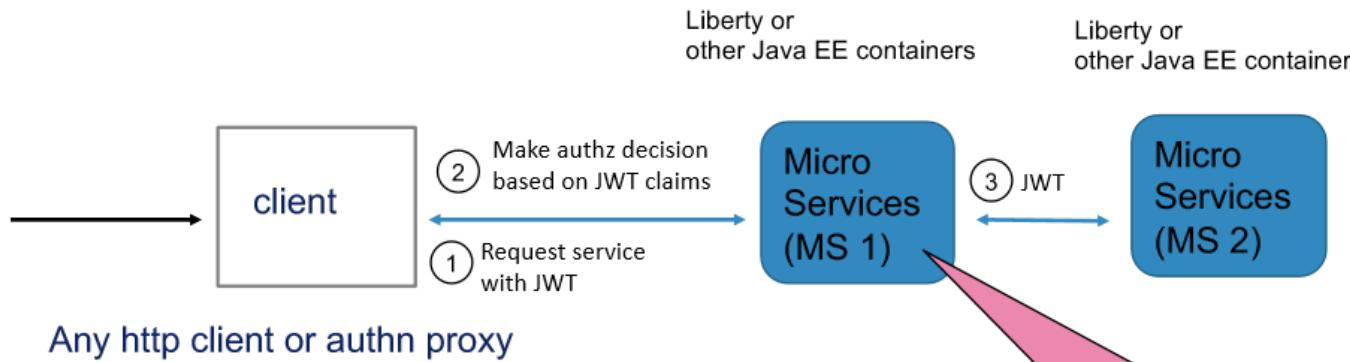
@Gauge(unit = MetricUnits.NONE,
       name = "inventorySizeGauge",
       absolute = true,
       description = "Number of systems in the inventory")
public int getTotal() {
    return systems.size();
}
```

```
# TYPE application_inventoryProcessingTime_seconds summary
# HELP application_inventoryProcessingTime_seconds Time needed to process the inventory
application_inventoryProcessingTime_seconds_count{method="list"} 2
application_inventoryProcessingTime_seconds{method="list",quantile="0.5"} 2.2185000000000002E-5
...
# TYPE application_inventoryAccessCount_total counter
# HELP application_inventoryAccessCount_total Number of times the list of systems method is requested
application_inventoryAccessCount_total 1

# TYPE application_inventorySizeGauge gauge
# HELP application_inventorySizeGauge Number of systems in the inventory
application_inventorySizeGauge 1
```

MicroProfile: JWT Propagation features

Config	Fault Tolerance	Health Check	Health Metrics	JWT Propagation
externalize configuration to improve portability	build robust behavior to cope with unexpected failures	common format to determine service availability	common REST endpoints for monitoring service health	interoperable authentication and role-based access control mpJwt-1.0



1. Client has JWT, and uses it to request service over http header
2. Service verifies JWT & creates JsonWebToken & subject
3. Service authorizes request with JsonWebToken

Communication between services

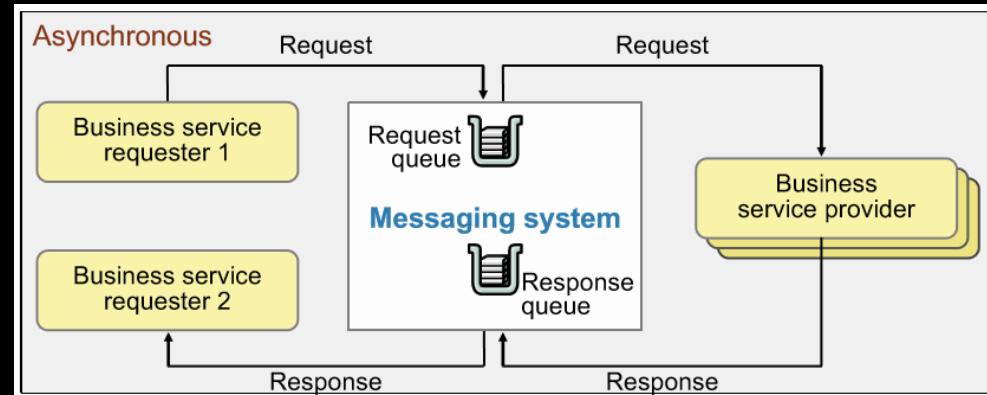
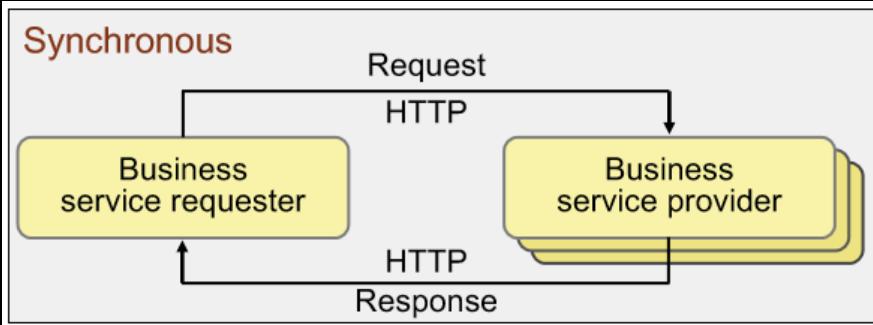
- Communication between services should be language-neutral
 - Services can be written in different languages
- Most often REST synchronous protocol
 - JSON, HTTP
 - Network/firewall friendly
- For asynchronous protocol – Event driven apps
 - Messaging system, typically MQ or Kafka
 - JSON message payloads
 - P2P or Pub/Sub

Asynchronous communication

Asynchronous communication can make microservices more robust

- Better decoupling
- Requester doesn't have to block while provider runs
- Different requester instance can handle response
- Messaging system holds action and result

Be aware of additional latency and unpredicted response time



Configuration for microservices

Make configuration part of your DevOps process

- Treat infrastructure as code, software-defined data center
- All application deployment must be automated
- Start small with simple scripts and build from there
- All configuration properties externalized as environment variables

Refactoring to Microservices

Evolving to microservices

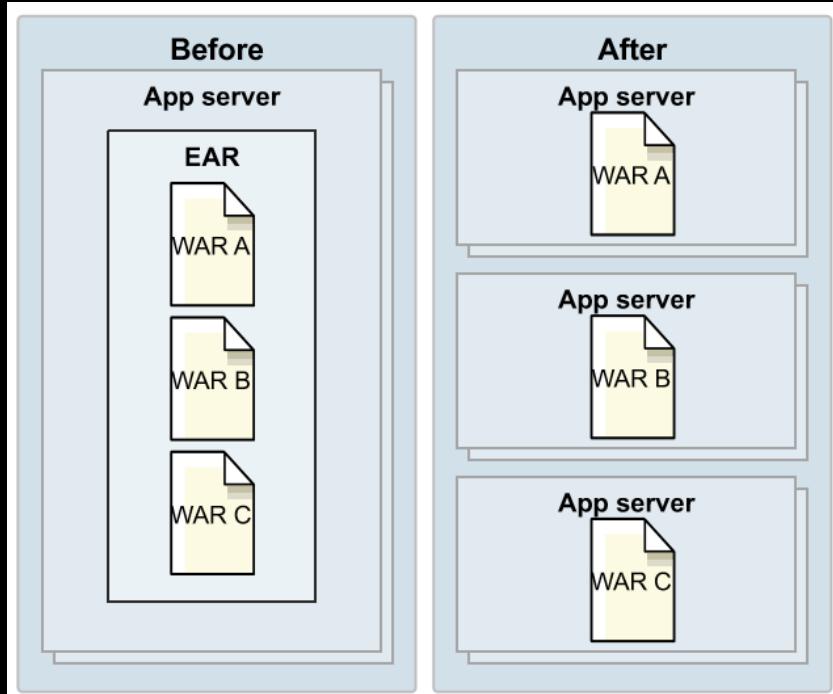
- Agile development: What's the simplest thing that could possibly work?
 - A monolith is simpler
- Start with a monolith
 - Make it modular, and plan that modules can become microservices
 - Each module should be a vertical slice, a mini-app with its own data
- Start with a minimally viable product (MVP)
 - MVP is a module – improvements to existing features go in the same module
 - Implement new functional tasks in new modules
 - When a module implements more than one task, refactor into separate modules
- Separate modules into microservices as needed
 - Multiple teams want to work independently
 - Monolithic code base is becoming unwieldy to maintain
 - Modules need to deploy, scale, or fail independently

Refactoring an existing application

- Places to refactor
 - Each REST service is potentially a microservice
 - Each SOAP web service or EJB is potentially a microservice
 - Especially stateless session beans
 - Redesign function-oriented interfaces to asset-oriented interfaces
 - Make the interface RESTful, such as using command objects
 - Use domain-driven design to discover your assets, which might be microservices
- Refactor the data for the microservices code
 - Each microservice manages its own data
 - A set of tables that is only used by one module
 - A set of tables that is only used at a particular time, like a daily status summary
 - Refactor the tables so that each table is used by only one module
 - Optimize for query time, not storage efficiency
- Consider runtime modernization
 - Check if monolith can be run as whole on modern runtime e.g. WebSphere Liberty

Refactoring an existing application

- Split up EAR files
 - Divide along functional lines
 - Package independent WAR files
 - Might require minor code changes
- Deploy WAR files separately
 - Its own Liberty for Java instant runtime
 - Its own Liberty Docker container
- Build, deploy, and manage separately
 - Use independent DevOps pipelines for each WAR file
 - Scale each separately and manage independently



Questions/Discussions?